

# report

June 8, 2025

## 1 Project Description

### 1.1 I. Project Goal

The primary goal of this project is to predict the percentage change in the price of Crude Oil futures two minutes following the release of the Energy Information Agency's (EIA) Weekly Petroleum Status Report (WPSR).

### 1.2 II. Data Overview - Quinn Kelly

This project utilizes several key datasets to achieve its objective:

1. **EIA WPSR Forecast Data (`InvestingcomEIA.csv`)**
  - This dataset comprises historical records of EIA WPSR releases (reported in EST). Key variables include:
    - **Release Date and Time:** The date and time (EST) of the EIA report release.
    - **Actual Change in Crude Oil Inventories:** The officially reported weekly change in U.S. commercial crude oil inventories, in millions of barrels (figures are suffixed with 'M').
    - **Market Forecasts:** Pre-release consensus forecasts for the change in crude oil inventories (in millions of barrels), sourced from a third-party provider.
    - **Previous Period's Actual:** The actual inventory change from the preceding reporting period.
2. **EIA WPSR Crude Oil Net Imports (`net_import.csv`)**
  - This dataset contains historical weekly U.S. crude oil net import volumes. Key variables are:
    - **Date:** The Friday marking the end of the week for which net imports were recorded.
    - **Net Import:** The average daily net import volume of crude oil for the week, in thousands of barrels per day.
3. **EIA WPSR Weekly U.S. Field Production of Crude Oil (`weekly_prod.csv`)**
  - This dataset provides historical weekly data on U.S. field production of crude oil. Key variables include:
    - **Date:** The Friday marking the end of the week for which production was recorded.
    - **Production:** The average daily U.S. field production of crude oil for the week, in thousands of barrels per day.
4. **Minute-Resolution WTI Crude Oil Price Data (`c1-1m.csv`)**
  - This dataset provides high-frequency, minute-by-minute price and volume data for WTI crude oil futures (data timestamped in GMT-6/CST). Key variables include:
    - **Date and Time:** The date and time (GMT-6/CST) of each minute bar.

- **Open, High, Low, Close (OHLC) Prices:** The open, high, low, and closing prices for each minute interval.
- **Volume:** The trading volume within each minute interval.

**Data Sourcing Note:** All EIA datasets (`InvestingcomEIA.csv`, `net_import.csv`, `weekly_prod.csv`) are publicly available online. The minute-resolution WTI crude oil price data (`c1-1m.csv`) was obtained from a commercial online market data provider.

### 1.3 III. Data Structure and Instances - Evan

- **Instance Definition:** Each instance in the final dataset used for machine learning represents a single weekly EIA WPSR release event.
- **Total Instances:** The dataset comprises **731** such instances.
- **Instance Composition:** Each instance is a consolidation of:
  - OHLCV (Open, High, Low, Close, Volume) price data for WTI crude oil futures, covering 60 minutes before and 2 minutes after the WPSR release.
  - Relevant statistics from the corresponding WPSR report, such as:
    - \* Change in crude oil inventories (supply)
    - \* Net imports
    - \* Production figures.

### 1.4 IV. Target Variable - Terence

The target variable we aim to predict is the **percentage change in the price of WTI crude oil futures, calculated two minutes after the WPSR release time relative to the price at release time.**

### 1.5 V. Features for Prediction - Samyak

The features used to predict the target variable include:

- **Pre-Release Price Data:** Sixty minutes of minute-by-minute OHLC price data for WTI crude oil futures immediately preceding the WPSR release.
- **WPSR Report Statistics:**
  - Actual change in crude oil inventories (supply)
  - Net imports
  - Production figures
- **Release Time Price:** The exact price of WTI crude oil futures at the moment of the WPSR release.

Problem Setup

## 2 Problem Setup - Mia

### 2.0.1 A. Data Cleaning/Pre-processing

Prior to model development, the raw data from each source required several cleaning and pre-processing steps to address inconsistencies and prepare it for analysis:

- **Timezone Standardization:**

- The datetime information in the EIA WPSR forecast data (reported in EST) and the Crude Oil Futures price data (reported in CST) were from different timezones. These were standardized to a single, consistent timezone (UTC) to ensure accurate temporal alignment of events.
- **Handling Missing Price Data (Zero-Volume Minutes):**
  - The Crude Oil Futures price data did not explicitly record minutes with zero trading activity. To create a continuous time series, these missing minute intervals were identified and imputed using forward-fill from the last known price.
- **Alignment of WPSR Data Indices:**
  - When integrating the various EIA WPSR datasets, an indexing discrepancy was identified. The WPSR supply data (from `InvestingcomEIA.csv`) was indexed by the report's release date. In contrast, the net import (`net_import.csv`) and production (`weekly_prod.csv`) datasets were indexed by the final Friday of the week to which the data pertained. These datasets were carefully aligned based on the relevant reporting week to ensure that all WPSR statistics corresponded to the correct release event.

## 2.0.2 B. Investigating Market Reaction to WPSR Release - Evan and Mia

Our initial approach for this project was driven by the fundamental hypothesis that the information contained within the EIA's Weekly Petroleum Status Report (WPSR) is rapidly absorbed by the market upon its release. We posited that the crude oil market subsequently self-corrects its price to reflect this new information. Consequently, our first step was to investigate the speed of this information absorption and to determine if the report induced statistically significant changes in price behavior. To validate our hypothesis, we conducted the following exploratory analyses:

### 1. Trading Volume Analysis:

- We first examined trading volume around the WPSR release time. As illustrated in the plot, there is a clear and significant spike in trading volume precisely at the time of the report's release. This observation strongly suggests that the market is indeed reacting to the information and actively trading based on it.

```
[1]: #run this cell to produce Trading Volume Plot also preprocessing in this cell
import pandas as pd
from datetime import datetime
import datetime
import pytz
import re
import xgboost as xgb
from sklearn.model_selection import TimeSeriesSplit, ParameterGrid
import numpy as np
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
import plotly.graph_objects as go

min_WTI = pd.read_csv(
    'data/cl-1m.csv',
    sep=';',
    header=None,
```

```

    names=['Date', 'Time', 'Open', 'High', 'Low', 'Close', 'Volume']
)

icom_eia_forecasts = pd.read_csv('InvestingcomEIA.csv')

def standardize_eia_datetime(row):
    """Standardizes datetime for EIA forecast"""
    release_date = pd.to_datetime(row['Release Date'], format='%d-%b-%y',
    ↪errors='coerce').date()
    release_time = pd.to_datetime(row['Time'], format='%H:%M', errors='coerce').
    ↪time()

    if pd.isna(release_date) or pd.isna(release_time):
        print('nan')
        return pd.NaT

    combined_datetime = pd.Timestamp.combine(release_date, release_time)
    return combined_datetime

icom_eia_forecasts['Release_Datetime'] = icom_eia_forecasts.
    ↪apply(standardize_eia_datetime, axis=1)

eastern_tz = pytz.timezone('US/Eastern')
icom_eia_forecasts['Release_Datetime_EST'] =
    ↪icom_eia_forecasts['Release_Datetime'].dt.tz_localize(eastern_tz,
    ↪ambiguous='infer', nonexistent='shift_forward')

chicago_tz = pytz.timezone('America/Chicago')
icom_eia_forecasts['Release_Datetime_CST'] =
    ↪icom_eia_forecasts['Release_Datetime_EST'].dt.tz_convert(chicago_tz)

def standardize_wti_datetime(row):
    """Standardizes datetime for min_WTI dataframe."""
    date_str = row['Date']
    time_str = row['Time']

    combined_datetime = pd.to_datetime(date_str + ' ' + time_str, format='%d/%m/
    ↪%Y %H:%M:%S', errors='coerce')

    if pd.isna(combined_datetime):
        print('nan')
        return pd.NaT
    return combined_datetime

min_WTI['Datetime'] = min_WTI.apply(standardize_wti_datetime, axis=1)

```

```

min_WTI['Datetime_CST'] = min_WTI['Datetime'].dt.tz_localize(chicago_tz,
↳ambiguous='infer', nonexistent='shift_forward')

icom_eia_forecasts = icom_eia_forecasts.set_index('Release_Datetime_CST').
↳sort_index()
min_WTI = min_WTI.set_index('Datetime_CST').sort_index()

min_WTI = min_WTI[~min_WTI.index.duplicated(keep='first')]

# filling in timejumps with 0 activity
min_WTI = min_WTI.resample('min').asfreq()
min_WTI['Close'] = min_WTI['Close'].ffill()
min_WTI['Open'] = min_WTI['Open'].fillna(min_WTI['Close'])
min_WTI['High'] = min_WTI['High'].fillna(min_WTI['Close'])
min_WTI['Low'] = min_WTI['Low'].fillna(min_WTI['Close'])
min_WTI['Volume'] = min_WTI['Volume'].fillna(0)
min_WTI['Date'] = min_WTI['Date'].ffill()

min_WTI = min_WTI.reset_index()

min_WTI['Datetime_CST'] = pd.to_datetime(min_WTI['Datetime_CST'])
min_WTI['Time'] = min_WTI['Datetime_CST'].dt.strftime('%H:%M:%S')
min_WTI['Datetime'] = min_WTI['Datetime_CST'].dt.strftime('%Y-%m-%d %H:%M:%S')
percentage_price_changes_1min_wti = []
previous_close_price = None

for index, row in min_WTI.iterrows():
    current_close_price = row['Close']
    if previous_close_price is not None and previous_close_price != 0:
        percentage_change = ((current_close_price - previous_close_price) /
↳previous_close_price) * 100.0
        percentage_price_changes_1min_wti.append(percentage_change)
    else:
        percentage_price_changes_1min_wti.append(float('nan'))
    previous_close_price = current_close_price

percentage_price_changes_1min_wti_series = pd.
↳Series(percentage_price_changes_1min_wti, index=min_WTI.index)
min_WTI['Percent_Change'] = percentage_price_changes_1min_wti_series
min_WTI.set_index('Datetime_CST', inplace=True)

```

```

def get_price_windows(eia_release_times, price_data, window_minutes_before=60,
    ↪window_minutes_after=60):
    """
    Extracts price data windows around EIA report release times.

    Args:
        eia_release_times (pd.DatetimeIndex): Index of icom_eia_forecasts
        ↪(release datetimes).
        price_data (pd.DataFrame): min_res_OIH dataframe with Datetime index.
        window_minutes_before (int): Minutes to include before release time.
        window_minutes_after (int): Minutes to include after release time.

    Returns:
        pd.DataFrame: A DataFrame containing price data for all events, within
        ↪the specified windows.
        Returns an empty DataFrame if no data is found within any
        ↪window.
    """
    price_windows_list = []

    for release_time in eia_release_times:
        start_time = release_time - pd.Timedelta(minutes=window_minutes_before)

        end_time = release_time + pd.Timedelta(minutes=window_minutes_after)

        window_data = price_data.loc[start_time:end_time].copy()

        if not window_data.empty:
            window_data['Release_Datetime'] = release_time
            price_windows_list.append(window_data)

    if price_windows_list:
        price_windows_df = pd.concat(price_windows_list)
        return price_windows_df
    else:
        return pd.DataFrame()

price_window_60min = get_price_windows(icom_eia_forecasts.index, min_WTI,
    ↪window_minutes_before=60, window_minutes_after=60)

price_window_60min = price_window_60min.reset_index()

price_window_60min['Time_to_Release_Minutes'] =
    ↪(price_window_60min['Datetime_CST'] -
    ↪price_window_60min['Release_Datetime']).dt.total_seconds() / 60

```

```

price_window_60min.set_index('Datetime_CST', inplace=True)
price_window_60min.sort_index(inplace=True)
average_volume_by_time = price_window_60min.
    ↪groupby('Time_to_Release_Minutes')['Volume'].mean().reset_index()

fig_volume_scatter_release = go.Figure(data=[go.Scatter(
    x=average_volume_by_time['Time_to_Release_Minutes'],
    y=average_volume_by_time['Volume'],
    mode='markers',
    marker=dict(size=8),
    text=average_volume_by_time['Volume'],
    hovertemplate="Time to Release: %{x:.0f} minutes<br>Average Volume: %{y:.
    ↪0f}<extra></extra>"
)])

fig_volume_scatter_release.update_layout(
    title='Average Trading Volume Around EIA WPSR Release',
    xaxis_title='Minutes Relative to EIA Report Release',
    yaxis_title='Average Volume',
    xaxis=dict(
        tickvals=[-60, -45, -30, -15, 0, 15, 30, 45, 60],
        ticktext=['-60', '-45', '-30', '-15', 'Release', '+15', '+30', '+45',
    ↪'+60']
    ),
    hovermode="closest"
)

fig_volume_scatter_release.show()

```

## 2. Price Behavior Analysis:

- Next, we investigated changes in price behavior surrounding the release. This was approached in two ways:
  - **Visual Inspection via Heatmaps:** We generated two heatmaps to compare minute-by-minute percentage price changes. In these heatmaps, each square's hue represents the frequency of specific price change magnitudes. One heatmap depicts price changes on days with a WPSR release, while the other shows price changes on non-release days. A visual comparison clearly indicates a distinct shift in price behavior at the time of the release, which gradually reverts to its typical pattern afterward.

```

[2]: #run this cell to produce the two heatmaps
time_intervals_release = sorted(price_window_60min['Time_to_Release_Minutes'].
    ↪unique())
price_change_bins = np.linspace(-1, 1, num=41)
price_change_labels = [f'{bin_val:.2f}%' for bin_val in price_change_bins]

```

```

price_window_60min['Price_Change_Bin'] = pd.
    ↳cut(price_window_60min['Percent_Change'], bins=price_change_bins,
    ↳labels=price_change_labels[:-1], include_lowest=True)

heatmap_data_release = price_window_60min.groupby(['Price_Change_Bin',
    ↳'Time_to_Release_Minutes'], observed=False).size().unstack(fill_value=0)

heatmap_data_release = heatmap_data_release.
    ↳reindex(columns=time_intervals_release, fill_value=0)
heatmap_data_release = heatmap_data_release.reindex(index=price_change_labels[:
    ↳-1], fill_value=0)

fig_heatmap_release = go.Figure(data=go.Heatmap(
    z=heatmap_data_release.values,
    x=heatmap_data_release.columns,
    y=heatmap_data_release.index,
    colorscale='Viridis',
    colorbar=dict(title='Frequency')
))

fig_heatmap_release.update_layout(
    title='Frequency of Percentage Price Changes Around EIA WPSR Release',
    xaxis_title='Time Relative to Release (Minutes)',
    yaxis_title='Percentage Price Change Bins',
    yaxis=dict(autorange="reversed"),
    xaxis=dict(tickvals=time_intervals_release[:,5], ticktext=[int(x) for x in
    ↳time_intervals_release[:,5]])
)

eia_release_dates = icom_eia_forecasts.reset_index()['Release_Datetime_CST'].dt.
    ↳date
min_wti_dates = min_WTI.reset_index()['Datetime_CST'].dt.date
mask_non_eia_days = ~min_wti_dates.isin(eia_release_dates)

min_WTI_no_eia_days = min_WTI.reset_index()[mask_non_eia_days].copy()

wti_with_eia_releases = pd.merge_asof(
    left=min_WTI_no_eia_days.reset_index(),
    right=icom_eia_forecasts.reset_index(),
    left_on='Datetime_CST',
    right_on='Release_Datetime_CST',
    direction='nearest',
    tolerance=pd.Timedelta('2D')
)

```



```

wti_with_eia_releases = wti_with_eia_releases.set_index('Datetime_CST')

wti_with_eia_releases.dropna(subset=['Release_Datetime_CST'], inplace=True)

def time_difference(time1, time2):
    """
    Calculates the time difference in seconds between two datetime.time objects.

    Args:
        time1: The first datetime.time object.
        time2: The second datetime.time object.

    Returns:
        The time difference in minutes as a float.
    """
    dummy_date = datetime.date(1, 1, 1)
    datetime1 = datetime.datetime.combine(dummy_date, time1)
    datetime2 = datetime.datetime.combine(dummy_date, time2)

    time_delta = datetime2 - datetime1

    return time_delta.total_seconds() / 60

time_diff = [time_difference(t1, t2) for t1, t2 in zip(wti_with_eia_releases.
    ↪reset_index()['Release_Datetime_CST'].dt.time, wti_with_eia_releases.
    ↪reset_index()['Datetime_CST'].dt.time)]

time_diff_series = pd.Series(time_diff)
wti_with_eia_releases.reset_index(inplace=True)
wti_with_eia_releases['Time_till_Release'] = (time_diff_series)
wti_with_eia_releases.set_index('Datetime_CST', inplace=True)

non_release_price_windows_60min = wti_with_eia_releases.copy().
    ↪loc[wti_with_eia_releases['Time_till_Release'].abs() <= 60.0]
non_release_price_windows_60min.reset_index(inplace=True)
non_release_price_windows_60min.drop(columns=['Datetime_CST', 'Open', 'High',
    ↪'Low', 'Close', 'Release_Datetime_CST', 'Release_Datetime_EST'],
    ↪inplace=True)

time_intervals_non_release =
    ↪sorted(non_release_price_windows_60min['Time_till_Release'].unique())

non_release_price_windows_60min['Price_Change_Bin'] = pd.
    ↪cut(non_release_price_windows_60min['Percent_Change'],
    ↪bins=price_change_bins, labels=price_change_labels[:-1], include_lowest=True)

```

```

heatmap_data_non_release = non_release_price_windows_60min.
    ↳groupby(['Price_Change_Bin', 'Time_till_Release'], observed=False).size().
    ↳unstack(fill_value=0)

heatmap_data_non_release = heatmap_data_non_release.
    ↳reindex(columns=time_intervals_non_release, fill_value=0)
heatmap_data_non_release = heatmap_data_non_release.
    ↳reindex(index=price_change_labels[:-1], fill_value=0)

heatmap_non_release_values = heatmap_data_non_release.values

min_val = np.min(heatmap_non_release_values)
max_val = np.max(heatmap_non_release_values)

if max_val > min_val:
    heatmap_non_release_values_scaled_linear = (heatmap_non_release_values -
    ↳min_val) / (max_val - min_val) * 100.0
else:
    heatmap_non_release_values_scaled_linear = np.
    ↳zeros_like(heatmap_non_release_values)

fig_heatmap_non_release = go.Figure(data=go.Heatmap(
    z=heatmap_non_release_values_scaled_linear,
    x=heatmap_data_non_release.columns,
    y=heatmap_data_non_release.index,
    colorscale='Viridis',
    colorbar=dict(title='Frequency')
))

fig_heatmap_non_release.update_layout(
    title='Frequency of Percentage Price Changes Around EIA WPSR Release On_
    ↳Non-Release Days',
    xaxis_title='Time Relative to Release (Minutes)',
    yaxis_title='Percentage Price Change Bins',
    xaxis=dict(tickvals=time_intervals_non_release[::5], ticktext=[int(x) for x_
    ↳in time_intervals_non_release[::5]])
)

fig_heatmap_release.show()
fig_heatmap_non_release.show()

```

- – **Statistical Verification via Kolmogorov-Smirnov Test:** To statistically confirm the observed changes in price behavior, we performed a two-sample Kolmogorov-Smirnov (KS) test. This test compared the distributions of minute-by-minute price changes immediately following the WPSR release on “release days” versus comparable times on

“non-release days.” The results, indicate that the price distributions/behavior during the first two minutes post-release are statistically significantly different from those on non-release days and from subsequent minutes on release days. *(Note: For this test, a p-value threshold of 0.1 was chosen instead of the more conventional 0.05 due to the limited number of WPSR release events available for comparison.)*

```
[3]: #run this cell to produce the KS Test polot
from scipy import stats
price_change_bins_labels = heatmap_data_non_release.index
price_change_bins_midpoints = []
for label in price_change_bins_labels:
    lower_bound_str = label.split('%')[0]
    midpoint = float(lower_bound_str) / 100.0
    price_change_bins_midpoints.append(midpoint)
price_change_bins_midpoints = np.array(price_change_bins_midpoints)

minutes_to_test = pd.Series([float(m) for m in range(0, 16)])

ks_results = []

for minute in minutes_to_test:
    if minute not in heatmap_data_non_release.columns or minute not in ↵
    ↪heatmap_data_release.columns:
        print(f"Minute {minute} not found in both heatmaps. Skipping.")
        continue

    sample_non_release = []
    frequencies_non_release = heatmap_data_non_release[minute]
    for i, freq in enumerate(frequencies_non_release):
        sample_non_release.extend([price_change_bins_midpoints[i]] * freq)
    sample_non_release = np.array(sample_non_release)

    sample_release = []
    frequencies_release = heatmap_data_release[minute]
    for i, freq in enumerate(frequencies_release):
        sample_release.extend([price_change_bins_midpoints[i]] * freq)
    sample_release = np.array(sample_release)

    if sample_non_release.size > 0 and sample_release.size > 0:
        ks_statistic, p_value = stats.ks_2samp(sample_non_release, ↵
        ↪sample_release)
        ks_results.append({
            'Minute': float(minute),
            'KS Statistic': ks_statistic,
            'P-value': p_value
        })
    else:
```

```

        ks_results.append({
            'Minute': float(minute),
            'KS Statistic': np.nan,
            'P-value': np.nan,
            'Warning': 'One or both samples are empty, KS test not performed.'
        })

ks_results_df = pd.DataFrame(ks_results)

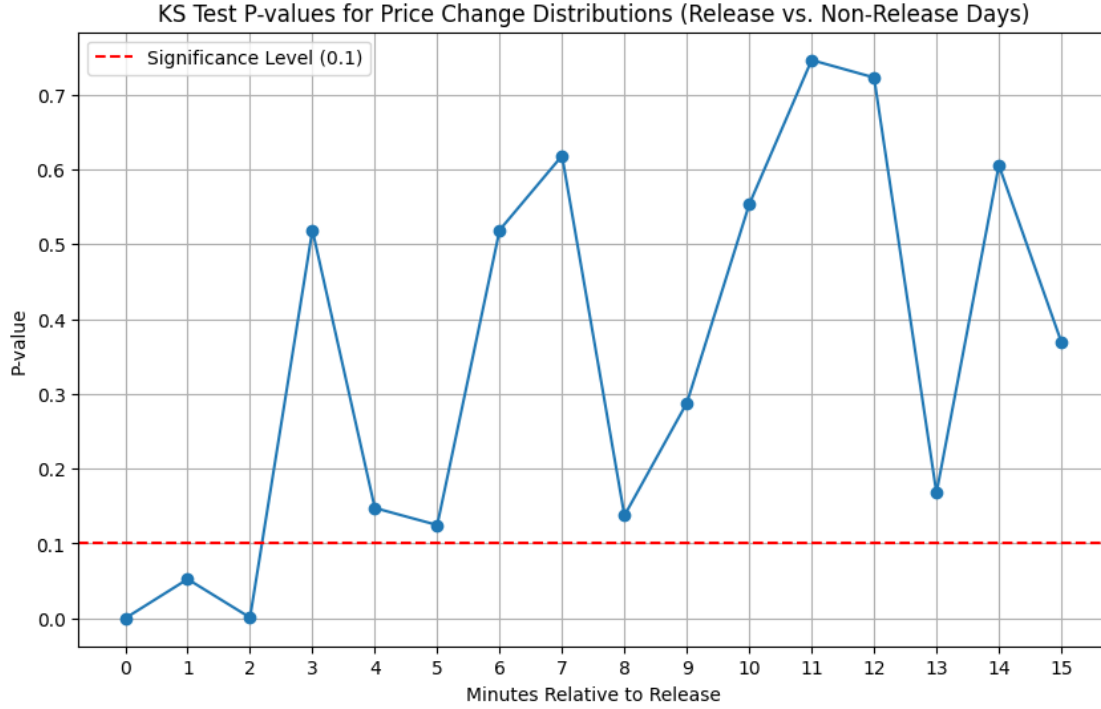
def create_ks_test_p_value_plot(ks_results_df):
    """
    Generates a line plot of KS test P-values for price change distributions.

    Args:
        ks_results_df: DataFrame containing 'Minute' and 'P-value' columns.

    Returns:
        matplotlib.figure.Figure: The generated plot figure.
    """
    line_plot_p_value_KS_test = plt.figure(figsize=(10, 6))
    plt.plot(ks_results_df['Minute'], ks_results_df['P-value'], marker='o',
    ↪linestyle='-')
    plt.axhline(0.1, color='r', linestyle='--', label='Significance Level (0.
    ↪1)')
    plt.title('KS Test P-values for Price Change Distributions (Release vs.
    ↪Non-Release Days)')
    plt.xlabel('Minutes Relative to Release')
    plt.ylabel('P-value')
    plt.xticks(ks_results_df['Minute'])
    plt.legend()
    plt.grid(True)
    return line_plot_p_value_KS_test

line_plot_p_value_KS_test = create_ks_test_p_value_plot(ks_results_df)

```



### 3 Algorithms

Having established that the initial two minutes of trading activity post-WPSR release represent the most relevant period for price movement, our focus shifted to predicting the percentage price change within this window using machine learning techniques.

Our group implemented and evaluated the following five regression algorithms:

- **Linear Regression** (Mia)
- **Random Forest Regression** (Quinn)
- **Extreme Gradient Boosting (XGBoost) Regression** (Evan)
- **Deep Neural Network (DNN)** (Terence)
- **K-Nearest Neighbors (KNN) Regression** (Samyak)

#### 3.1 Target Variable Justification: Percentage Price Change - Everyone

For this project, the target variable we aim to predict is the **percentage change in the WTI crude oil futures price, calculated two minutes after the WPSR release relative to the price at the moment of release**. This choice is deliberate and offers several key advantages for modeling financial time series, particularly in the context of event-driven price movements:

##### 1. Normalization and Comparability:

- Percentage change inherently normalizes the price movement by the prevailing price level. A \$0.50 price move has a vastly different significance if crude oil is trading at \$20/barrel versus \$100/barrel. By using percentage change, we ensure that the target

variable represents a relative magnitude of change, making predictions comparable across different time periods and varying absolute price levels. This is crucial as our dataset spans a considerable time, during which oil prices have fluctuated significantly.

## 2. Stationarity Properties:

- Raw price series in finance are often non-stationary (e.g., they can exhibit trends, random walks, or unit roots). Non-stationary data can lead to spurious correlations and unreliable models. While not always perfectly stationary, log returns (which percentage changes closely approximate for small changes) tend to exhibit more stationary behavior than absolute price levels or even absolute price differences. This makes them more amenable to standard machine learning regression techniques.

## 3. Addressing Heteroscedasticity (to some extent):

- Financial time series often exhibit volatility clustering (periods of high volatility followed by periods of low volatility). The variance of absolute price changes can be dependent on the price level (higher prices often mean larger absolute fluctuations). Percentage changes can help mitigate this effect, as the volatility of returns is often more stable.

```
[4]: #pre processing
valid_releases = price_window_60min.
    ↳loc[price_window_60min['Time_to_Release_Minutes'] == 2]
price_window_60min = price_window_60min[price_window_60min['Date']].
    ↳isin(valid_releases['Date'])]
price_window_60min =
    ↳price_window_60min[price_window_60min['Time_to_Release_Minutes'] <= 2]

def pivot_market_data(df, x_minutes_before, cols_to_pivot=None):
    """
    pivot into wide formate where minutes to release from [-x, 2] are kept as
    ↳columns for each feature.
    """
    if cols_to_pivot is None:
        cols_to_pivot = ['Open', 'High', 'Low', 'Close', 'Volume']

    df_filtered = df[(df['Time_to_Release_Minutes'] <= 2) &
    ↳(df['Time_to_Release_Minutes'] >= -x_minutes_before)].copy()

    df_long = df_filtered.melt(
        id_vars=['Datetime', 'Time_to_Release_Minutes', 'Release_Datetime'],
        value_vars=cols_to_pivot,
        var_name='Feature',
        value_name='Value'
    )
    df_long['Feature_min'] = df_long['Feature'] + '_t' +
    ↳df_long['Time_to_Release_Minutes'].astype(int).astype(str)

    df_wide = df_long.pivot_table(
        index='Release_Datetime',
        columns='Feature_min',
```

```

        values='Value'
    ).reset_index()
    cols_to_drop = [col for col in df_wide.columns if 't2' in col and col != 'Close_t2']
    df_wide.drop(columns=cols_to_drop, inplace=True)
    return df_wide

df_wide = pivot_market_data(price_window_60min, x_minutes_before=60)

def get_time_offset(col_name):
    match = re.search(r't(-?\d+)$', col_name)
    if match: return int(match.group(1))
    return float('inf')

sorted_market_cols = sorted(df_wide.columns, key=get_time_offset)
df_wide = df_wide[sorted_market_cols]
df_wide.set_index('Release_Datetime', inplace=True)
df_wide['Price_Change'] = df_wide['Close_t2'] - df_wide['Open_t0']
#df_wide is now sorted

df_wide.reset_index(inplace=True)
df_wide['Release Date'] = df_wide['Release_Datetime'].dt.date
df_wide = df_wide.drop(columns=['Release_Datetime']).set_index('Release Date')

weekly_supply=pd.read_csv("InvestingcomEIA.csv")
weekly_supply['Release Date'] = pd.to_datetime(weekly_supply['Release Date'],
    format="%d-%b-%y").dt.strftime("%Y-%m-%d")
weekly_supply.set_index('Release Date', inplace=True)

weekly_prod=pd.read_csv("weekly_prod.csv", header=2)

weekly_import=pd.read_csv("net_import.csv")

df_wide = df_wide.sort_index()
df_wide.reset_index(inplace=True)
df_wide['Pct_Change'] = (df_wide['Close_t2'] - df_wide['Open_t0']) / df_wide['Open_t0'] * 100
weekly_supply = weekly_supply.sort_index()
weekly_supply.reset_index(inplace=True)

df_wide['Release Date'] = pd.to_datetime(df_wide['Release Date'])
weekly_supply['Release Date'] = pd.to_datetime(weekly_supply['Release Date'])
price_supply = pd.merge(df_wide, weekly_supply, how='left', on='Release Date')

weekly_prod['Date'] = pd.to_datetime(weekly_prod['Date'])

```

```

weekly_import['Date'] = pd.to_datetime(weekly_import['Date'])

prod_import = pd.merge(weekly_import, weekly_prod, how='left', on='Date').
    ↪dropna()

mapping = pd.read_csv('FXmappings.csv')
split_data = mapping['Date'].str.split('(', expand=True)

mapping['ReleaseDate_str'] = split_data[0].str.strip()

mapping['end_date'] = split_data[1].str.rstrip(')').str.strip()

mapping['ReleaseDate'] = pd.to_datetime(mapping['ReleaseDate_str'], format='%m/
    ↪%d/%Y')
mapping.drop(columns=['ReleaseDate_str', 'Date'], inplace=True)

release_year = mapping['ReleaseDate'].dt.year

full_end_date = mapping['end_date'] + ' ' + release_year.astype(str)

mapping['end_date'] = pd.to_datetime(full_end_date, format='%b %d %Y')

mapping['day_before'] = mapping['ReleaseDate'] - pd.Timedelta(days=1)

dow = mapping['day_before'].dt.dayofweek

days_to_subtract = (dow + 3) % 7

mapping['end_date'] = mapping['day_before'] - pd.to_timedelta(days_to_subtract,
    ↪unit='D')

mapping.drop(columns=['day_before'], inplace=True)
mapping['end_date'] = pd.to_datetime(mapping['end_date'])
mapping = mapping.set_index('end_date').sort_index()
mapping.reset_index(inplace=True)

prod_import_mapped = pd.merge(mapping, prod_import, how='left',
    ↪right_on='Date', left_on='end_date')
prod_import_mapped.drop(columns=['end_date', 'Date'], inplace=True)

full_data = pd.merge(price_supply, prod_import_mapped, how='left',
    ↪left_on='Release Date', right_on='ReleaseDate')

full_data.drop(columns=['Time', 'ReleaseDate'], inplace=True)
full_data['Weekly Net Import'] = full_data['Weekly Net Import'] * 7
full_data['Weekly Production'] = full_data['Weekly U.S. Field Production of
    ↪Crude Oil (Thousand Barrels per Day)'] * 7

```



```

full_data = full_data.drop(columns=['Weekly U.S. Field Production of Crude Oil ↵
↵(Thousand Barrels per Day)'])
full_data = full_data.dropna()

def convert_to_numeric(value):
    multiplier = 1_000_000
    numeric_part = value[:-1]

    try:
        return float(numeric_part) * multiplier
    except ValueError:
        return np.nan

for col in ['Actual', 'Forecast', 'Previous']:

    full_data[col] = full_data[col].apply(convert_to_numeric)

for col in ['Weekly Net Import', 'Weekly Production']:
    full_data[col] = full_data[col].apply(lambda x: x * 1000)

```

## 3.2 Linear Regression (Mia Callahan)

### 3.2.1 Design Justifications

The Linear Regression model, including its Lasso-regularized variant, was implemented with several key design decisions aimed at establishing a robust baseline, understanding linear relationships, and managing potential overfitting given the high dimensionality of the pre-release market data.

#### 3.2.2 1. Choice of Linear Regression (and Lasso)

- **Decision:** Employ Linear Regression, subsequently augmented with Lasso (L1) regularization.
- **Justification:**
  - **Baseline Performance:** Linear Regression serves as a fundamental benchmark. Its performance provides a clear baseline against which more complex models (Random Forest, XGBoost, DNN) can be quantitatively compared. This helps in assessing whether the added complexity of other models translates into a significant improvement in predictive power for this specific financial prediction task.
  - **Interpretability:** Standard Linear Regression coefficients (especially after feature scaling) offer direct insights into the linear relationship (strength and direction) between each feature and the percentage price change. Lasso, by shrinking some coefficients to zero, further enhances interpretability by performing feature selection.
  - **Simplicity and Speed:** Linear models are computationally inexpensive and fast to train, making them ideal for initial data exploration and for establishing a quick performance reference.

- **Identifying Linear Signals:** Despite the expectation of non-linearities in financial markets, this model helps to identify and quantify any strong, simple linear patterns that might be present in the feature set.

### 3.2.3 2. Data Splitting Strategy

- **Decision:** Utilize a standard `train_test_split` methodology to create training, validation, and test sets (resulting in an 80% train, 10% validation, and 10% test split based on the code structure). `random_state=42` is used for reproducibility.
- **Justification:**
  - **Simplicity for Baseline Model:** For establishing a baseline, a randomized train-validation-test split is straightforward to implement and provides a quick assessment of general predictive capability on unseen data.
  - **Validation Set for Hyperparameter Tuning:** The creation of a separate validation set is crucial for tuning the `alpha` parameter of the Lasso regression without contaminating the final test set.

### 3.2.4 3. Feature Scaling

- **Decision:** Apply `StandardScaler` to the input features before training the Linear Regression and Lasso models.
- **Justification:**
  - **Equal Feature Contribution:** Linear Regression models, and particularly regularized versions like Lasso, are sensitive to the scale of input features. Features with larger numerical ranges can disproportionately influence the model's coefficient fitting process. `StandardScaler` normalizes features by removing the mean and scaling to unit variance. This ensures all features are on a comparable scale, allowing the model to learn their importance based on their relationship with the target, not their magnitude.
  - **Improved Regularization Performance:** For Lasso regression, which applies an L1 penalty to the coefficients, feature scaling is important for fair penalization. Without scaling, features with larger values would naturally have larger coefficients, leading to them being penalized more heavily, irrespective of their true predictive power. Scaling helps the regularization process to be more effective and stable.
  - **Numerical Stability:** Can improve the numerical stability and convergence speed of the optimization algorithms used in fitting the models.

### 3.2.5 4. Lasso Regression for Feature Selection and Regularization

- **Decision:** Implement Lasso (L1 Regularization) Regression and tune its `alpha` hyperparameter using the validation set.
- **Justification:**
  - **Addressing High Dimensionality & Preventing Overfitting:** The feature set includes 60 minutes of OHLCV data, leading to a large number of predictors (300 market features + `Open_t0` + 4 fundamental features = 305 features before any interaction terms). Standard Linear Regression with many features is prone to overfitting the training data. Lasso adds an L1 penalty term to the loss function, which shrinks some less important coefficients towards zero, and can force some to be exactly zero. This penalizes model complexity and helps improve generalization to unseen data.

- **Automatic Feature Selection:** By setting irrelevant feature coefficients to zero, Lasso performs an implicit form of feature selection. This is particularly useful when dealing with many potentially correlated or noisy features, as it helps in identifying a more parsimonious and interpretable model.
- **Managing Multicollinearity:** While not its primary function, Lasso can handle multicollinearity by tending to select one feature from a group of highly correlated features and shrinking the coefficients of others.
- **Alpha Hyperparameter Tuning:** The **alpha** parameter controls the strength of the L1 penalty. A small alpha leads to behavior similar to ordinary least squares, while a large alpha increases shrinkage, potentially leading to underfitting. The process of iterating through a range of **alpha** values and selecting the one that minimizes Mean Squared Error (MSE) on the validation set (as indicated by the code logic) is a sound method for finding an optimal trade-off between bias and variance.

### 3.2.6 5. Evaluation Metric

- **Decision:** Use Root Mean Squared Error (RMSE) to evaluate model performance on the test set.
- **Justification:**
  - **Standard Regression Metric:** RMSE is a widely accepted and standard metric for evaluating the accuracy of regression models.
  - **Penalizes Larger Errors More Heavily:** Due to the squaring of errors before averaging, RMSE gives a higher weight to larger errors. In financial predictions, large, unexpected errors can be particularly detrimental, making RMSE a suitable metric if the goal is to minimize such occurrences.

These design decisions collectively aim to build a well-regularized linear model that can serve as a solid baseline, provide insights into linear feature importance, and offer a fair comparison point for more advanced techniques applied in the project.

Running a linear regression acts as a baseline for our other models. We expect that the linear regression will not perform well, as the data likely has non-linear relationships. The bias-variance tradeoff likely will lead the linear regression to underfit the data as it is a very basic model. However, it will be interesting to see how it performs especially in comparison to the more complex models we will use later.

First, we plot each variable against our outcome variable to determine if linear relationships between the variable can be estimated. (Plotted last 20 features for readability follow comments to see every feature plotted)

```
[5]: #run to see feature linear plots

newy = df_wide['Pct_Change']
feature_cols = [col for col in df_wide.columns if '_t2' not in col and '_t_1'
↳not in col and '_t_0' not in col and col not in ['Close_t2',
↳'_Release_Datetime', 'Date', 'Pct_Change', 'Price_Change']] + ['Open_t0']
feature_cols = [col for col in feature_cols if col not in ['Unnamed: 0',
↳'_Release Date']]
ncols = 10
```


```

nrows = 2 #(-len(feature_cols) // ncols)
fig, axes = plt.subplots(nrows, ncols, figsize=(15, 2.5* nrows))
axes = axes.flatten()
X = df_wide[feature_cols]
y = df_wide['Close_t2']

import warnings
from sklearn.exceptions import ConvergenceWarning
warnings.filterwarnings('ignore', category=ConvergenceWarning)

for idx,col in enumerate(feature_cols[-20:]): # delete -20
    ax = axes[idx]
    xplot = X[col]
    ax.plot(xplot,newy, "b.")
    ax.set_xlabel(col, fontsize=10)
plt.tight_layout()
plt.show()

```



Unsurprisingly, this leads to a much less linear relationship between the data. While this will lead to worse performance of my model, the intended use of prediction models is to be able to profit off of stock market changes, so this is a more useful metric.

```

[8]: from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

n_splits_outer = 3
tscv_outer = TimeSeriesSplit(n_splits=n_splits_outer)

for train_val_indices, test_indices in tscv_outer.split(X):
    pass

```

```

X_train_val_temp, y_train_val_temp = X.iloc[train_val_indices], newy.
    ↪iloc[train_val_indices]
X_test, y_test = X.iloc[test_indices], newy.iloc[test_indices]

n_splits_inner = 3

tscv_inner = TimeSeriesSplit(n_splits=n_splits_inner)
for train_idx_inner, val_idx_inner in tscv_inner.split(X_train_val_temp):
    pass # Ensures we get the last set of indices
train_indices = train_idx_inner
val_indices = val_idx_inner

X_train, y_train = X_train_val_temp.iloc[train_indices], y_train_val_temp.
    ↪iloc[train_indices]
X_val, y_val = X_train_val_temp.iloc[val_indices], y_train_val_temp.
    ↪iloc[val_indices]

print("Data Split Sizes:")
print(f"Total samples: {len(X)}")
print(f"Training samples: {len(X_train)} (~{len(X_train)/len(X)*100:.1f}%)")
print(f"Validation samples: {len(X_val)} (~{len(X_val)/len(X)*100:.1f}%)")
print(f"Test samples: {len(X_test)} (~{len(X_test)/len(X)*100:.1f}%)")

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_val_scaled = scaler.transform(X_val)
X_test_scaled = scaler.transform(X_test)

np.random.seed(42)
linmodel = LinearRegression()
linmodel.fit(X_train_scaled, y_train)

print(f"Intercept: {linmodel.intercept_}")
y_val_pred = linmodel.predict(X_val_scaled)
val_rmse = np.sqrt(mean_squared_error(y_val, y_val_pred))
print(f"Validation RMSE: {val_rmse:.4f}")

# Predictions on the test set
y_test_pred = linmodel.predict(X_test_scaled)
test_rmse = np.sqrt(mean_squared_error(y_test, y_test_pred))
print(f"Test RMSE: {test_rmse:.4f}")

```

Data Split Sizes:  
Total samples: 731

Training samples: 412 (~56.4%)  
Validation samples: 137 (~18.7%)  
Test samples: 182 (~24.9%)  
Intercept: -0.02721175775627985  
Validation RMSE: 0.3799  
Test RMSE: 0.5521

Because there are so many features it is likely that the linear regression will overfit because it will try to fit all of the features, even if they are not relevant. After seeing performance, we will reduce the number of features to see if this improves out of sample performance.

The RMSE of the model on the test set is .55, which is much better than I expected. This could change for a larger data sample. To try to prevent any overfitting from the large number of features, I will now use lasso to select the most useful features. First I will cross validate to find the ideal alpha.

```
[11]: from sklearn.linear_model import Lasso
      mses = []
      best_mse = float('inf')
      best_alpha = None
      alphas = np.logspace(-10, 0, 100)
      for alpha in alphas:
          lasso = Lasso(alpha=alpha, random_state=42)
          lasso.fit(X_train_scaled, y_train)
          mse = mean_squared_error(y_val, lasso.predict(X_val_scaled))
          mses.append(mse)
          if mse < best_mse:
              best_mse = mse
              best_alpha = alpha
      print(f"Best alpha: {best_alpha}")
```

Best alpha: 0.09770099572992247

```
[12]: lasso_best = Lasso(alpha=best_alpha, random_state=42)
      lasso_best.fit(X_train_scaled, y_train)
      # see how it performs on out of sample data
      y_hat_lasso = lasso_best.predict(X_test_scaled)
      error_lasso = np.sqrt(np.mean((y_hat_lasso - y_test) ** 2))
      print(f"RMSE on test set with Lasso: {error_lasso:.2f}")
```

RMSE on test set with Lasso: 0.27

The RMSE is higher than that using the full set of features. This indicates that most features are useful in improving out of sample prediction.

### 3.3 Random Forest (Quinn Kelly)

#### 3.3.1 Why Random Forest Was Chosen for Implementation:

Random Forest is a robust and widely used ensemble learning method, and its selection for predicting financial metrics (like price changes based on news sentiment) can be justified for several

key reasons:

**1. Ability to Capture Non-Linear Relationships:**

- Financial markets are notoriously complex, and the relationship between input features (e.g., news sentiment, technical indicators, economic data) and the target variable (e.g., price change) is rarely linear. Random Forest, being an ensemble of decision trees, is inherently capable of modeling complex, non-linear interactions between features without requiring explicit transformation or specification of these interactions beforehand.

**2. Robustness to Overfitting (Compared to Single Decision Trees):**

- While individual decision trees can easily overfit the training data, Random Forest mitigates this risk through two main mechanisms:
  - **Bagging (Bootstrap Aggregating):** Each tree in the forest is trained on a random bootstrap sample of the training data.
  - **Feature Randomness:** At each split in a tree, only a random subset of features is considered.
- This combination creates diverse trees, and averaging their predictions tends to reduce variance and lead to better generalization on unseen data. The use of `ccp_alpha` for pruning further helps control complexity and prevent overfitting.

**3. Good Performance with Less Hyperparameter Tuning (Relatively):**

- Compared to some other advanced models (like Support Vector Machines or Neural Networks), Random Forests can often achieve good performance with less extensive hyperparameter tuning. They are often a strong “out-of-the-box” performer, making them a good choice for establishing a solid baseline or even as a primary model.

**4. Handling of Different Feature Types and Scales:**

- Random Forests can handle numerical features directly without requiring extensive preprocessing like feature scaling (though scaling can sometimes still be beneficial for other parts of a pipeline or for consistency). This simplifies the preprocessing pipeline.

**5. Implicit Feature Importance Estimation:**

- Random Forests provide a natural way to estimate the importance of each feature in making predictions (e.g., by measuring how much each feature contributes to reducing impurity across all trees). In a financial context, understanding which factors (e.g., sentiment from specific news categories, particular economic indicators) are most influential can be as valuable as the prediction itself.

**6. Robustness to Outliers and Noisy Data:**

- Financial data is often noisy and can contain outliers. The ensemble nature of Random Forest, where predictions are averaged across many trees, makes the model more robust to the influence of individual outliers or noise points in the training data compared to models that might be heavily skewed by such points.

**7. Good Choice as a Benchmark or Comparison Model:**

- Even if more complex models like Gradient Boosting (XGBoost, LightGBM) or Neural Networks are also being considered, Random Forest serves as an excellent benchmark. Its performance can provide a strong baseline against which other, potentially more complex or harder-to-tune, models can be compared. It represents a powerful, yet relatively understandable, class of non-linear models.

### 3.3.2 Design Justifications:

#### 3.3.3 1. Hyperparameters Grid and Tuning Strategy

- **Decision** - 'n\_estimators\_list': [200, 400, 600], 'ccp\_alpha\_list': [10\*\*-2, 10\*\*-3, 10\*\*-5, 10\*\*-7]:
  - **n\_estimators**: This parameter defines the number of trees in the forest. A larger number of trees generally improves the model's performance and stability by reducing variance, but at an increased computational cost. The list [200, 400, 600] explores a range of forest sizes to balance predictive power and efficiency.
  - **ccp\_alpha (Minimal Cost-Complexity Pruning)**: This parameter controls the complexity of the individual trees within the forest. Non-negative values of **ccp\_alpha** act as a threshold for pruning; trees (or subtrees) that do not offer sufficient improvement in terms of impurity reduction (scaled by **ccp\_alpha**) are pruned. Larger **ccp\_alpha** values lead to more aggressive pruning and simpler trees, which can help prevent overfitting. The selected range [10\*\*-2, 10\*\*-3, 10\*\*-5, 10\*\*-7] allows for testing various degrees of pruning, from relatively strong to very mild.

#### 3.3.4 2. Time Series Cross-Validation and Data Splitting

- **Decision** - **tscv = TimeSeriesSplit(n\_splits=n\_split)** (e.g., **n\_split=3**):
  - **Justification**: Standard k-fold cross-validation is inappropriate for time-series data because it can lead to lookahead bias (training on future data to predict the past). **TimeSeriesSplit** ensures that the training set always precedes the validation/test set. In each split, the training set grows, and the test set is the next chronological block of data, mimicking how a model would be deployed in real-time. **n\_split=3** creates 3 such chronological train/test splits.

```
[13]: from sklearn.model_selection import TimeSeriesSplit
from sklearn.ensemble import RandomForestRegressor
import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import mean_squared_error

def time_series_valid_test(X, y, n_split, valid_or_test, n_estimators_list,
    ↪ccp_alpha_list, optimal_par=None):
    np.random.seed(42)
    tscv = TimeSeriesSplit(n_splits=n_split)
    rf_rmse = []
    rmseSmall = []
    param_combinations = [] # param combo storagae
    i = 0 # fold counter
    for train_index, test_index in tscv.split(X):
        i += 1
        # Break test set into 50% validation set, 50% test set
        break_test_ind = int(test_index[0] + 0.5*(test_index[-1]-test_index[0]))
        valid_index = np.array(list(range(test_index[0],break_test_ind)))
        test_index = np.array(list(range(break_test_ind,test_index[-1])))
```



```

    # Split data into training, validation, and test sets
    X_train, X_valid, X_test = X.iloc[train_index], X.iloc[valid_index], X.
↪iloc[test_index]
    y_train, y_valid, y_test = y.iloc[train_index], y.iloc[valid_index], y.
↪iloc[test_index]

    # Tuning
    if valid_or_test == "valid":
        rf_rmse_fold = []
        for n_estimators in n_estimators_list:
            for ccp_alpha in ccp_alpha_list:
                model_rf = RandomForestRegressor(random_state=42,
                                                n_estimators=n_estimators,
                                                ccp_alpha=ccp_alpha,
                                                n_jobs=-1)

                X_train_sample = X_train.sample(frac=0.1, random_state=42)
                y_train_sample = y_train.loc[X_train_sample.index]

                model_rf.fit(X_train_sample, y_train_sample.to_numpy())

                y_val_rf = model_rf.predict(X_valid)
                fold_rmse = np.sqrt(mean_squared_error(y_valid, y_val_rf))
                rf_rmse_fold.append(fold_rmse)
                param_combinations.append((n_estimators, ccp_alpha))

        rf_rmse.append(np.mean(rf_rmse_fold))

    if valid_or_test == "test":
        model_rf = RandomForestRegressor(random_state=42,
                                        n_estimators=optimal_par[0],
                                        ccp_alpha=optimal_par[1],
                                        n_jobs=-1)

        model_rf.fit(X_train, y_train.to_numpy())
        y_test_rf = model_rf.predict(X_test)
        test_rmse = np.sqrt(mean_squared_error(y_test, y_test_rf))
        rf_rmse.append(test_rmse)

    # Plot the prediction for the last CV fold
    if i == n_split:
        plt.plot(range(len(X_test)), y_test_rf, label="Prediction")
        plt.plot(range(len(X_test)), y_test, label="True Value")
        plt.legend(loc="upper left")
        plt.show()

```

```

# Return results
if valid_or_test == "valid":
    min_rmse_idx = np.argmin(rf_rmse) # Find index of lowest RMSE
    optimal_params = param_combinations[min_rmse_idx] # Retrieve
    ↪corresponding parameters
    return rf_rmse, optimal_params
if valid_or_test == "test":
    rf_rmse = np.mean(rf_rmse)
    rmseSmallAns = np.mean(rmseSmall)
    return y_test_rf, y_test

```

```

[14]: # Hyperparameter combinations
n_estimators_list = [200, 400, 600]
ccp_alpha_list = [10**-2, 10**-3, 10**-5, 10**-7]

rf_rmse, optimal_params = time_series_valid_test(X, newy, n_split = 3,
    ↪valid_or_test = "valid",
                                                    n_estimators_list =
    ↪n_estimators_list,
                                                    ccp_alpha_list =
    ↪ccp_alpha_list)
print("Optimal Parameters:", optimal_params)
print("Validation RMSEs (Validation):", rf_rmse)

```


Optimal Parameters: (200, 1e-05)  
 Validation RMSEs (Validation): [0.7044190597545149, 0.336229695292928,  
 0.28911024341411684]

```

[15]: y_test_rf, y_test_actual = time_series_valid_test(X, newy, n_split = 3,
    ↪valid_or_test = "test",
                                                    n_estimators_list =
    ↪n_estimators_list,
                                                    ccp_alpha_list = ccp_alpha_list,
                                                    optimal_par = optimal_params)

print("RMSE on test set at final fold:", np.sqrt(mean_squared_error(y_test_rf,
    ↪y_test_actual)))

```



RMSE on test set at final fold: 0.2523032045361738

c:\Users\evans\AppData\Local\Programs\Python\Python310\lib\site-packages\numpy\core\fromnumeric.py:3504: RuntimeWarning:

Mean of empty slice.

c:\Users\evans\AppData\Local\Programs\Python\Python310\lib\site-packages\numpy\core\\_methods.py:129: RuntimeWarning:

invalid value encountered in scalar divide

Interpretations: Random forest was able to generate a better approximation than LASSO / linear regression. This was performed by taking a time series train-test split and then using k-fold cross validation on a validation set to tune the following parameters to generate the best fitting random forest model: `n_estimators` and `ccp_alpha` from the following respective choices: [200,400, 600] and [10-2,10-3, 10-5,10-7]. The goal was to lower the root mean squared error between the two respective data sets, where the RMSE was derived from predictions the model generated at the final fold. For our purposes that was 3. When this was performed on the smaller set, we were able to receive a rmse of 0.427. Once the model was returned to the larger, more dense dataset, we were able to calibrate a model to achieve a rmse of 0.29. As we would expect, the denser data were able to produce a lower rmse indicating the additional columns in the data matrix were positive

additions to the study.

### 3.4 Extreme Gradient Boost (XGBoost) Regression (Evan Sun)

Reasons I chose to implement XGBoost:

**1. Superior Predictive Performance and Robustness:**

- XGBoost is a gradient boosting algorithm renowned for its high predictive accuracy. It often outperforms other algorithms in machine learning competitions and real-world applications (I've lost too many kaggle competitions to XGBoost), including finance.
- The ensemble nature of XGBoost, where multiple decision trees are built sequentially with each new tree correcting the errors of the previous ones, leads to a robust model that is less prone to individual tree weaknesses.

**2. Effective Handling of Non-Linearity and Feature Interactions:**

- Financial markets, and specifically price responses to news events like the EIA WPSR, rarely exhibit simple linear relationships. Price changes can be influenced by a complex interplay of various factors (e.g., the magnitude of the surprise in the report, prevailing market sentiment, existing inventory levels, recent price trends).
- XGBoost, being a tree-based model, excels at capturing these non-linear relationships and high-order interactions between features without requiring explicit feature engineering for these interactions (e.g., creating polynomial terms).

**3. Built-in Regularization to Prevent Overfitting:**

- Financial data is notoriously noisy, and overfitting is a significant risk. XGBoost incorporates both L1 (Lasso) and L2 (Ridge) regularization terms in its objective function.
- This regularization helps to penalize model complexity, shrink less important feature weights, and thus improve the model's ability to generalize to unseen data – crucial for predicting future price movements after a WPSR release. Your hyperparameter grid explicitly tunes `reg_alpha` (L1) and `reg_lambda` (L2), leveraging this strength.

**4. Computational Efficiency and Scalability:**

- XGBoost is designed for speed and efficiency. It implements parallel processing for tree construction and cache-aware access patterns.

<!--also the name is cool lol-->

```
[16]: #create training and test split
column_names = []
variable_names = ['Open', 'High', 'Low', 'Close', 'Volume']
for var_name in variable_names:
    for min in range(60):
        column_names.append(f'{var_name}_t{min-60}')

column_names.append('Open_t0')

new_cols = ['Actual', 'Forecast', 'Previous', 'Weekly Net Import', 'Weekly_
↪Production']
for col_name in new_cols:
    column_names.append(col_name)

n_total_events = len(full_data)
```

```

test_set_size_ratio = 0.2

test_set_split_index = int(n_total_events * (1 - test_set_size_ratio))

# Data for hyperparameter tuning and training the final model
X_tuning_full = full_data.iloc[:test_set_split_index][column_names]
y_tuning_full = full_data.iloc[:test_set_split_index]['Pct_Change']

# Final hold-out test set (never seen during tuning)
X_test_holdout = full_data.iloc[test_set_split_index:][column_names]
y_test_holdout = full_data.iloc[test_set_split_index:]['Pct_Change']

print(f"Full dataset size: {len(df_wide)}")
print(f"Tuning dataset size: {len(X_tuning_full)}")
print(f"Hold-out test dataset size: {len(X_test_holdout)}")

```

Full dataset size: 731

Tuning dataset size: 586

Hold-out test dataset size: 147

### 3.4.1 Design Justifications:

Some of the design decisions made might not be conventional and thus can be explained as follows:

### 3.4.2 1. Hyperparameters Grid

- **Decision - 'n\_estimators': [1000]:** I chose the max value here since I'm using early stopping so I can let early stopping determine the optimal number of trees dynamically for each hyperparameter combination and fold. The `actual_n_estimators` is then recorded and averaged.

### 3.4.3 2. Time Series Cross-Validation

- **Decision - 'tscv = TimeSeriesSplit(n\_splits=3):** I'm using a time series split because for financial time series data (like price changes), it's crucial that the validation set always comes *after* the training set to prevent lookahead bias and realistically simulate how the model would be used in practice (training on past data to predict future data).

### 3.4.4 3. Early Stopping

- **Decision - 'early\_stopping\_rounds\_val = 10:** I chose to use early stopping for 2 reasons:
  - Prevents the model from fitting the training data too closely and losing generalization ability.
  - saves computation time by not necessarily training for the full `n_estimators` if optimal performance is achieved earlier.

### 3.4.5 4. Sample Weighting

- **Decision - 'epsilon = 1e-6:**
  - **Decision:** Add a small constant to target values before taking absolute value for weights.

- **Justification:** This is to prevent issues if any `y_tuning_full` or `y_test_holdout` values are exactly zero. `np.abs(0)` is 0, which would result in a zero weight, effectively ignoring that sample. `epsilon` ensures all samples have at least a tiny positive weight.
- **Decision - `weights_tuning_full = np.abs(y_tuning_full) + epsilon`:** I noticed that the 2-minute price changes were often very close to 0 and as a result I decided to implement sample weighting so that the model is incentivized to predict more important/larger price changes more accurately. This way the model doesn't default to just predicting 0 for most events and basically being useless.

```
[17]: #train and tune model
hyperparam_grid = {
    'n_estimators': [1000],
    'learning_rate': [0.01, 0.05, 0.1],
    'max_depth': [6, 7, 8],
    'subsample': [0.85, 0.9, 0.95],
    'colsample_bytree': [0.85, 0.9, 0.95],
    'reg_alpha': [0.1, 0.3, 0.5],
    'reg_lambda': [3, 4, 5],
}

param_list = list(ParameterGrid(hyperparam_grid))

tscv = TimeSeriesSplit(n_splits=3)
results = []
early_stopping_rounds_val = 10

epsilon = 1e-6
weights_tuning_full = np.abs(y_tuning_full) + epsilon
weights_test_holdout = np.abs(y_test_holdout) + epsilon

print(f"Starting hyperparameter tuning with TimeSeriesSplit. Total combinations:
↳ {len(param_list)}")
for fold, (train_index, val_index) in enumerate(tscv.split(X_tuning_full)):
    print(f"\nFold {fold+1}/{tscv.get_n_splits()}")
    X_train_fold, y_train_fold = X_tuning_full.iloc[train_index], y_tuning_full.
↳iloc[train_index]
    X_val_fold, y_val_fold = X_tuning_full.iloc[val_index], y_tuning_full.
↳iloc[val_index]
    weights_train_fold = weights_tuning_full.iloc[train_index]
    weights_val_fold = weights_tuning_full.iloc[val_index]

    print(f"  Training on {len(X_train_fold)} instances, validating on
↳ {len(X_val_fold)} instances.")
```

```

for i, params in enumerate(param_list):
    if (i + 1) % 10 == 0 or i == 0 or i == len(param_list) - 1 :
        print(f"    Combination {i+1}/{len(param_list)}: {params}")

    model = xgb.XGBRegressor(
        objective='reg:squarederror',
        random_state=42,
        early_stopping_rounds=early_stopping_rounds_val,
        n_jobs=-1,
        **params
    )

    model.fit(X_train_fold, y_train_fold,
              sample_weight=weights_train_fold,
              eval_set=[(X_val_fold, y_val_fold)],
              sample_weight_eval_set=[weights_val_fold],
              verbose=False)

    y_pred_val = model.predict(X_val_fold)
    val_rmse = np.sqrt(mean_squared_error(y_val_fold, y_pred_val,
    ↪sample_weight=weights_val_fold))

    result_entry = params.copy()
    result_entry['fold'] = fold + 1
    result_entry['val_rmse'] = val_rmse
    result_entry['actual_n_estimators'] = model.best_iteration
    results.append(result_entry)

results_df = pd.DataFrame(results)

grouping_cols = [col for col in hyperparam_grid.keys() if col != 'n_estimators']
average_metrics_across_folds = results_df.groupby(grouping_cols).agg(
    avg_val_rmse=('val_rmse', 'mean'),
    avg_actual_n_estimators=('actual_n_estimators', 'mean')
).reset_index()

best_params_row = average_metrics_across_folds.
    ↪loc[average_metrics_across_folds['avg_val_rmse'].idxmin()]

best_hyperparams = best_params_row.to_dict()
best_hyperparams['max_depth'] = int(best_hyperparams.get('max_depth'))
best_avg_rmse = best_hyperparams.pop('avg_val_rmse')
best_actual_n_estimators = int(round(best_hyperparams.
    ↪pop('avg_actual_n_estimators')))

```

```

print("\nBest hyperparameters found via TimeSeriesSplit CV:")
for param_name, param_val in best_hyperparams.items():
    print(f" {param_name}: {param_val}")
print(f" (Best actual n_estimators found by early stopping, averaged):
↪{best_actual_n_estimators}")
print(f" Average Validation RMSE across folds: {best_avg_rmse:.4f}")

```

Starting hyperparameter tuning with TimeSeriesSplit. Total combinations: 729

Fold 1/3

Training on 148 instances, validating on 146 instances.

Combination 1/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 10/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 20/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 30/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 40/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 50/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 60/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 70/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 80/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 90/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 100/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 110/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.9}



Combination 120/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 130/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 140/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 150/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 160/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 170/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 180/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 190/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 200/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 210/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 220/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 230/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 240/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 250/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 260/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 270/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 280/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 290/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 300/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 310/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 320/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 330/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 340/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 350/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 360/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 370/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 380/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 390/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 400/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 410/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 420/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 430/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 440/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 450/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 460/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 470/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 480/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 490/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 500/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 510/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 520/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 530/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 540/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 550/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 560/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 570/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 580/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 590/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 600/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 610/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 620/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 630/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 640/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 650/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 660/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 670/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 680/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 690/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 700/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 710/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 720/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 729/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.95}

Fold 2/3

Training on 294 instances, validating on 146 instances.

Combination 1/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 10/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 20/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 30/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 40/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 50/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 60/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 70/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 80/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 90/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 100/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 110/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 120/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 130/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 140/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 150/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 160/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 170/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 180/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 190/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 200/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 210/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 220/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 230/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 240/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 250/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 260/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 270/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 280/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 290/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 300/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 310/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 320/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 330/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 340/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 350/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 360/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 370/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 380/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 390/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 400/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 410/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 420/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 430/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 440/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 450/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 460/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 470/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 480/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 490/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 500/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 510/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 520/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 530/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 540/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 550/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 560/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 570/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 580/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 590/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 600/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 610/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 620/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 630/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 640/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.85}



Combination 650/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 660/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 670/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 680/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 690/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 700/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 710/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 720/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 729/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.95}

Fold 3/3

Training on 440 instances, validating on 146 instances.

Combination 1/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 10/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 20/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 30/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 40/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 50/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 60/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 70/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 80/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 90/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 100/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 110/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 120/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 130/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 140/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 150/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 160/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 170/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 180/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 190/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 200/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 210/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 220/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 230/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 240/729: {'colsample\_bytree': 0.85, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 250/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 260/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 270/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 280/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 290/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 300/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 310/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 320/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 330/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 340/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 350/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 360/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 370/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 380/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 390/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 400/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 410/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 420/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 430/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 440/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 450/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 460/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 470/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 480/729: {'colsample\_bytree': 0.9, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 490/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 500/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 510/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 520/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 530/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 540/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 550/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 560/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.01, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 570/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 580/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 590/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 600/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 610/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 620/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 630/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 640/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.05, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 3, 'subsample': 0.85}

Combination 650/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 3, 'subsample': 0.9}

Combination 660/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 3, 'subsample': 0.95}

Combination 670/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 6, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 4, 'subsample': 0.85}

Combination 680/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 4, 'subsample': 0.9}

Combination 690/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 4, 'subsample': 0.95}

Combination 700/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 7, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.85}

Combination 710/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.1, 'reg\_lambda': 5, 'subsample': 0.9}

Combination 720/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.3, 'reg\_lambda': 5, 'subsample': 0.95}

Combination 729/729: {'colsample\_bytree': 0.95, 'learning\_rate': 0.1, 'max\_depth': 8, 'n\_estimators': 1000, 'reg\_alpha': 0.5, 'reg\_lambda': 5, 'subsample': 0.95}

Best hyperparameters found via TimeSeriesSplit CV:

learning\_rate: 0.1

max\_depth: 8

subsample: 0.95

colsample\_bytree: 0.85

reg\_alpha: 0.1

reg\_lambda: 5.0

(Best actual n\_estimators found by early stopping, averaged): 7

Average Validation RMSE across folds: 0.8257

```
[18]: #evaluate model on test set
print("\n--- Training Final Model with Best Hyperparameters ---")
params_for_final_model = best_hyperparams.copy()

final_model = xgb.XGBRegressor(
    objective='reg:squarederror',
    random_state=42,
    n_jobs=-1,
    n_estimators=best_actual_n_estimators,
    **params_for_final_model
)

print(f"Training final model on {len(X_tuning_full)} instances.")
final_model.fit(X_tuning_full, y_tuning_full,
                sample_weight=weights_tuning_full,
                verbose=False)

print("\n--- Evaluating Final Model on Hold-Out Test Set ---")
y_pred_test_holdout = final_model.predict(X_test_holdout)

test_rmse = np.sqrt(mean_squared_error(y_test_holdout, y_pred_test_holdout,
                                       sample_weight=weights_test_holdout))

print(f"Final Model Performance on Hold-Out Test Set:")
```

```

print(f"  Test RMSE: {test_rmse:.6f}")

baseline_pred_zero = np.zeros_like(y_test_holdout)
baseline_rmse_zero = np.sqrt(mean_squared_error(y_test_holdout,
↪baseline_pred_zero,
↪sample_weight=weights_test_holdout))

print(f"\nBaseline Model (Predict Zero Change) Performance on Hold-Out Test Set:
↪")
print(f"  Baseline RMSE: {baseline_rmse_zero:.6f}")

try:
    feature_importances = final_model.feature_importances_
    importance_df = pd.DataFrame({'feature': X_tuning_full.columns,
↪'importance': feature_importances})
    importance_df = importance_df.sort_values(by='importance', ascending=False)
    print("\nTop 10 Feature Importances from Final Model:")
    print(importance_df.head(10))
except Exception as e:
    print(f"Could not retrieve feature importances: {e}")

```

--- Training Final Model with Best Hyperparameters ---  
 Training final model on 586 instances.

--- Evaluating Final Model on Hold-Out Test Set ---  
 Final Model Performance on Hold-Out Test Set:  
 Test RMSE: 0.393174

Baseline Model (Predict Zero Change) Performance on Hold-Out Test Set:  
 Baseline RMSE: 0.402660

Top 10 Feature Importances from Final Model:

	feature	importance
261	Volume_t-39	0.206890
248	Volume_t-52	0.066634
0	Open_t-60	0.059081
295	Volume_t-5	0.055731
251	Volume_t-49	0.053814
271	Volume_t-29	0.051934
297	Volume_t-3	0.035970
4	Open_t-56	0.034913
301	Actual	0.032565
244	Volume_t-56	0.029560

**Evaluation of Results:**

## 1. Model Performance vs. Baseline:

- **Final Model Test RMSE: 0.393174**
- **Baseline Model (Predict Zero Change) RMSE: 0.402660**
- **Interpretation:** My XGBoost model has a lower RMSE on the hold-out test set than the naive baseline model (which simply predicts no price change). This indicates that the model has learned some predictive patterns from the features and is providing more accurate predictions than simply assuming the price will remain static over the next 2 minutes.
- **Magnitude of Improvement:** The improvement in my model's RMSE is approximately  $(0.009486 / 0.402660) * 100\% = 2.35\%$  lower than the baseline.

## 2. Top 10 Feature Importances:


- 261 Volume\_t-39 (0.206890): This is by far the **most dominant feature**, contributing over 20% to the model's predictive power (based on this importance type, likely 'gain' or 'weight'). It suggests that the trading volume 39 minutes *before the WPSR release* is a very strong indicator of the subsequent price change. This could reflect anticipatory trading or positioning by informed market participants.
- 248 Volume\_t-52 (0.066634)
- 0 Open\_t-60 (0.059081)
- 295 Volume\_t-5 (0.055731)
- 251 Volume\_t-49 (0.053814)
- 271 Volume\_t-29 (0.051934)
- '297 Volume\_t-3 (0.0359

```
[19]: plt.figure(figsize=(15, 6))
time_indices = X_test_holdout.index

true_changes_flat = y_test_holdout.flatten() if hasattr(y_test_holdout,
↳ 'flatten') else y_test_holdout
predicted_changes_flat = y_pred_test_holdout.flatten() if
↳ hasattr(y_pred_test_holdout, 'flatten') else y_pred_test_holdout
comparison_df = pd.DataFrame({
    'True Change': true_changes_flat,
    'Predicted Change': predicted_changes_flat
})
plt.plot(time_indices, comparison_df['True Change'], label='True Change',
↳ marker='.', linestyle='-', alpha=0.8)
plt.plot(time_indices, comparison_df['Predicted Change'], label='Predicted
↳ Change', marker='.', linestyle='-', alpha=0.8)
plt.xlabel("Sample Index / Time")
plt.ylabel("Change")
plt.title(f"True vs. Predicted Change Over Time")
plt.legend()
plt.grid(True)
plt.tight_layout() # Adjust layout to make room for labels
```



```
plt.show()
```



### 3.5 Deep Neural Network (Terence Chiu)

#### Overview

This model implements a Deep Neural Network (DNN) to predict percentage changes in price following economic data releases. The model takes into account both historical price data and weekly economic indicators to make predictions.

#### Model Architecture

- **Input Layer:** 7 features
  - 4 price-related features: Close prices at t-60, t-40, t-20, and Open price at t0
  - 3 weekly economic indicators: Weekly Production, Weekly Net Import, and Actual Supply
- **Hidden Layers:**
  - First layer: 32 neurons with ReLU activation
  - Second layer: 8 neurons with ReLU activation
  - Dropout rate of 0.2 after each hidden layer
- **Output Layer:** Single neuron for percentage change prediction
- **Learning Rate:** 0.000187 (optimized through hyperparameter tuning, check `terenceTuning-Model.py`)

#### Data Processing

- **Feature Scaling:** All input features are standardized using `StandardScaler` - **Target Variable:** Percentage change in price 2 minutes after release - **Train-Test Split:** 80-20 time-based split to maintain temporal order - **Sample Weighting:** Implemented to give more importance to larger price movements

#### Training Process

- **Optimizer:** Adam optimizer - Adapts learning rates for each parameter individually
- **Loss Function:** Mean Squared Error (MSE) - - Aligns with our goal of minimizing prediction errors since it penalizes larger errors

- **Metrics:**
  - Custom RMSE metric - matches evaluation metric used in financial analysis, and more sensitive to outliers than MAE
  - Mean Absolute Error (MAE) - robust for extreme data
- **Early Stopping:** - Prevent overfitting
  - Monitors validation RMSE
  - Patience of 10 epochs
  - Restores best weights

#### Model Performance

- **Evaluation Metric:** Weighted RMSE - **Baseline Comparison:** Model performance is compared against a baseline of predicting zero change - **Feature Importance:** Analyzed using gradient-based importance scores to understand feature contributions

#### Key Innovations

1. **Sample Weighting:** Implementation of sample weights to handle the imbalance between small and large price movements 2. **Custom RMSE Metric:** Development of a custom RMSE metric for better model evaluation 3. **Feature Importance Analysis:** Gradient-based analysis to understand feature contributions 4. **Time-Based Split:** Maintaining temporal order in train-test split to prevent look-ahead bias

#### Model Limitations

1. Relies on only four prices within 60 minutes prior to weekly report release 2. Does not look at whole market fluctuations 3. Performance depends on the quality and timeliness of the economic data releases

#### Future Improvements

1. Experiment with different network architectures, such as LSTM, CNN? 2. Incorporate additional features or market indicators 3. Implement ensemble methods with other models 4. Explore different loss functions for better handling of extreme price movements

#### Data Retrieval

```
[20]: import tensorflow as tf
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error

df = pd.read_csv("full_data.csv")

feature_cols = [col for col in df.columns if 'Close_t-60' in col or
↳ 'Close_t-40' in col or 'Close_t-20' in col or col == 'Release Date' or col
↳ == 'Actual' or col == 'Weekly Net Import' or col == 'Weekly Production' or
↳ col == 'Open_t0']

X_temp = df[feature_cols]
y_temp = (df['Close_t2'] - df['Close_t0'])/df['Close_t0']

prod_weekly = X_temp[['Release Date', 'Weekly Production']]
```

```

net_import_weekly = X_temp[['Release Date', 'Weekly Net Import']]
supply_weekly = X_temp[['Release Date', 'Actual']]
price_wide = X_temp[['Release Date', 'Close_t-60', 'Close_t-40', 'Close_t-20',
↪ 'Open_t0']]

```

Data Scaling - for price features, weekly features, and target price feature

```

[21]: price_scaler = StandardScaler()
      target_scaler = StandardScaler()

      price_features = price_wide[['Close_t-60', 'Close_t-40', 'Close_t-20',
↪ 'Open_t0']]

      # Scale the price features in the dataframe
      for col in ['Close_t-60', 'Close_t-40', 'Close_t-20', 'Open_t0']:
          price_features[col] = price_scaler.fit_transform(price_wide[col].values.
↪ reshape(-1, 1)).flatten()

      # Scale the target values in the dataframe
      y_temp = target_scaler.fit_transform(y_temp.values.reshape(-1, 1)).flatten()

      #Scaler for weekly data
      weekly_scaler = StandardScaler()

      weekly_production_scaled = weekly_scaler.fit_transform(prod_weekly['Weekly
↪ Production'].values.reshape(-1, 1)).flatten()
      weekly_import_scaled = weekly_scaler.fit_transform(net_import_weekly['Weekly
↪ Net Import'].values.reshape(-1, 1)).flatten()
      weekly_supply_scaled = weekly_scaler.fit_transform(supply_weekly['Actual'].
↪ values.reshape(-1, 1)).flatten()

      X = []
      y = []

      for idx, row in price_features.iterrows():
          # Target: price of future 2 minutes after release (already scaled)
          target_price = y_temp[idx]

          production_value = weekly_production_scaled[idx]
          import_value = weekly_import_scaled[idx]
          supply_value = weekly_supply_scaled[idx]

```

```

    row_data = [price_features['Close_t-60'].
↪values[idx],price_features['Close_t-40'].
↪values[idx],price_features['Close_t-20'].
↪values[idx],price_features['Open_t0'].
↪values[idx],production_value,import_value,supply_value]
    X.append(row_data)
    y.append(target_price)

X = np.array(X)
y = np.array(y)

```

C:\Users\evans\AppData\Local\Temp\ipykernel\_4368\1293597208.py:8:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\evans\AppData\Local\Temp\ipykernel\_4368\1293597208.py:8:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\evans\AppData\Local\Temp\ipykernel\_4368\1293597208.py:8:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\evans\AppData\Local\Temp\ipykernel\_4368\1293597208.py:8:  
SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

Importing Model from `terenceModel.py`. Model decision/architecture decided from `terenceTuneModel` using `keras_tuner`. Best hyperparameters found from tuning

1. Best configuration:
2. 2 layers: [32 units, 0.2 dropout], [8 units, 0.2 dropout]
3. Learning rate: 0.000187

```
[27]: from terenceModel import DNN
      from terenceTrainModel import plot_predictions

      # Time-based 80/20 split
      n = len(X)
      split_idx = int(n * 0.8)
      X_train, X_test = X[:split_idx], X[split_idx:]
      y_train, y_test = y[:split_idx], y[split_idx:]
```

We applied sample weights in training of model, and the calculation of RMSE, as (close to) 0 percent change was something that was likely to occur. This model will be shown to be outperforming the baseline model.

```
[28]: # Compute sample weights for training and test sets
      epsilon = 1e-6
      sample_weights_train = np.abs(y_train) + epsilon
      sample_weights_test = np.abs(y_test) + epsilon

      model = DNN()

      # Train model with sample weights
      trained_model, _ = model.train(X_train, y_train,
      ↪sample_weight=sample_weights_train)

      # Predict on test set
      y_pred = trained_model.predict(X_test).flatten()

      # Weighted RMSE for test set
      weighted_test_rmse = np.sqrt(mean_squared_error(y_test, y_pred,
      ↪sample_weight=sample_weights_test))
      print(f"Weighted Test RMSE: {weighted_test_rmse:.4f}")
```

Epoch 1/100

```
1/15 [=>...] - ETA: 5s - loss: 2.0089 - rmse: 1.1155 -
mae: 0.7994WARNING:tensorflow:`evaluate()` received a value for `sample_weight`,
but `weighted_metrics` were not provided. Did you mean to pass metrics to
`weighted_metrics` in `compile()`? If this is intentional you can pass
```

```

`weighted_metrics=[]` to `compile()` in order to silence this warning.
15/15 [=====] - 1s 10ms/step - loss: 2.2948 - rmse:
1.1892 - mae: 0.9012 - val_loss: 2.7153 - val_rmse: 1.0969 - val_mae: 0.8224
Epoch 2/100
 1/15 [=>...] - ETA: 0s - loss: 2.6603 - rmse: 1.2678 -
mae: 0.9407WARNING:tensorflow:`evaluate()` received a value for `sample_weight`,
but `weighted_metrics` were not provided. Did you mean to pass metrics to
`weighted_metrics` in `compile()`? If this is intentional you can pass
`weighted_metrics=[]` to `compile()` in order to silence this warning.
15/15 [=====] - 0s 3ms/step - loss: 2.3328 - rmse:
1.1523 - mae: 0.8569 - val_loss: 2.7205 - val_rmse: 1.0990 - val_mae: 0.8243
Epoch 3/100
 1/15 [=>...] - ETA: 0s - loss: 2.5490 - rmse: 1.2828 -
mae: 0.9646WARNING:tensorflow:`evaluate()` received a value for `sample_weight`,
but `weighted_metrics` were not provided. Did you mean to pass metrics to
`weighted_metrics` in `compile()`? If this is intentional you can pass
`weighted_metrics=[]` to `compile()` in order to silence this warning.
15/15 [=====] - 0s 3ms/step - loss: 2.5625 - rmse:
1.1926 - mae: 0.8799 - val_loss: 2.7206 - val_rmse: 1.1003 - val_mae: 0.8260
Epoch 4/100
 1/15 [=>...] - ETA: 0s - loss: 0.3866 - rmse: 0.9058 -
mae: 0.7062WARNING:tensorflow:`evaluate()` received a value for `sample_weight`,
but `weighted_metrics` were not provided. Did you mean to pass metrics to
`weighted_metrics` in `compile()`? If this is intentional you can pass
`weighted_metrics=[]` to `compile()` in order to silence this warning.
15/15 [=====] - 0s 3ms/step - loss: 2.3521 - rmse:
1.1239 - mae: 0.8239 - val_loss: 2.7209 - val_rmse: 1.1015 - val_mae: 0.8276
Epoch 5/100
 1/15 [=>...] - ETA: 0s - loss: 1.6602 - rmse: 1.1580 -
mae: 0.9357WARNING:tensorflow:`evaluate()` received a value for `sample_weight`,
but `weighted_metrics` were not provided. Did you mean to pass metrics to
`weighted_metrics` in `compile()`? If this is intentional you can pass
`weighted_metrics=[]` to `compile()` in order to silence this warning.
15/15 [=====] - 0s 3ms/step - loss: 2.2456 - rmse:
1.1162 - mae: 0.8280 - val_loss: 2.7232 - val_rmse: 1.1033 - val_mae: 0.8296
Epoch 6/100
 1/15 [=>...] - ETA: 0s - loss: 1.4675 - rmse: 1.0361 -
mae: 0.7686WARNING:tensorflow:`evaluate()` received a value for `sample_weight`,
but `weighted_metrics` were not provided. Did you mean to pass metrics to
`weighted_metrics` in `compile()`? If this is intentional you can pass
`weighted_metrics=[]` to `compile()` in order to silence this warning.
15/15 [=====] - 0s 3ms/step - loss: 2.3649 - rmse:
1.1030 - mae: 0.8059 - val_loss: 2.7223 - val_rmse: 1.1037 - val_mae: 0.8303
Epoch 7/100
 1/15 [=>...] - ETA: 0s - loss: 2.7301 - rmse: 1.2808 -
mae: 0.9858WARNING:tensorflow:`evaluate()` received a value for `sample_weight`,
but `weighted_metrics` were not provided. Did you mean to pass metrics to
`weighted_metrics` in `compile()`? If this is intentional you can pass

```

```

`weighted_metrics=[]` to `compile()` in order to silence this warning.
15/15 [=====] - 0s 3ms/step - loss: 2.3118 - rmse:
1.1368 - mae: 0.8456 - val_loss: 2.7259 - val_rmse: 1.1054 - val_mae: 0.8318
Epoch 8/100
 1/15 [=>...] - ETA: 0s - loss: 1.7810 - rmse: 1.2385 -
mae: 1.0111WARNING:tensorflow:`evaluate()` received a value for `sample_weight`,
but `weighted_metrics` were not provided. Did you mean to pass metrics to
`weighted_metrics` in `compile()`? If this is intentional you can pass
`weighted_metrics=[]` to `compile()` in order to silence this warning.
15/15 [=====] - 0s 3ms/step - loss: 2.2089 - rmse:
1.1218 - mae: 0.8273 - val_loss: 2.7338 - val_rmse: 1.1085 - val_mae: 0.8343
Epoch 9/100
 1/15 [=>...] - ETA: 0s - loss: 0.4815 - rmse: 0.8199 -
mae: 0.6109WARNING:tensorflow:`evaluate()` received a value for `sample_weight`,
but `weighted_metrics` were not provided. Did you mean to pass metrics to
`weighted_metrics` in `compile()`? If this is intentional you can pass
`weighted_metrics=[]` to `compile()` in order to silence this warning.
15/15 [=====] - 0s 3ms/step - loss: 2.5000 - rmse:
1.1350 - mae: 0.8088 - val_loss: 2.7363 - val_rmse: 1.1097 - val_mae: 0.8354
Epoch 10/100
 1/15 [=>...] - ETA: 0s - loss: 0.8956 - rmse: 0.9753 -
mae: 0.8358WARNING:tensorflow:`evaluate()` received a value for `sample_weight`,
but `weighted_metrics` were not provided. Did you mean to pass metrics to
`weighted_metrics` in `compile()`? If this is intentional you can pass
`weighted_metrics=[]` to `compile()` in order to silence this warning.
15/15 [=====] - 0s 3ms/step - loss: 2.1273 - rmse:
1.0391 - mae: 0.7576 - val_loss: 2.7353 - val_rmse: 1.1106 - val_mae: 0.8369
Epoch 11/100
 1/15 [=>...] - ETA: 0s - loss: 0.6303 - rmse: 0.7953 -
mae: 0.6372WARNING:tensorflow:`evaluate()` received a value for `sample_weight`,
but `weighted_metrics` were not provided. Did you mean to pass metrics to
`weighted_metrics` in `compile()`? If this is intentional you can pass
`weighted_metrics=[]` to `compile()` in order to silence this warning.
15/15 [=====] - 0s 3ms/step - loss: 2.1428 - rmse:
1.0697 - mae: 0.7902 - val_loss: 2.7372 - val_rmse: 1.1119 - val_mae: 0.8383
5/5 [=====] - 0s 2ms/step
Weighted Test RMSE: 1.5660

```

```

[29]: baseline_pred_zero = np.zeros_like(y_test)
baseline_rmse_zero = np.sqrt(mean_squared_error(y_test, baseline_pred_zero,
↪sample_weight=sample_weights_test))
print(f"Baseline (Zero Prediction) Weighted RMSE: {baseline_rmse_zero:.4f}")

if weighted_test_rmse < baseline_rmse_zero:
    print("Model outperforms baseline (predicting zero change).")
else:
    print("Model does NOT outperform baseline (predicting zero change).")


```

```
plot_predictions(y_pred, y_test)
```

Baseline (Zero Prediction) Weighted RMSE: 1.5964

Model outperforms baseline (predicting zero change).

Test RMSE: 1.1196



We compare with the weightless model with the same settings except that this model does not train with `sample_weights`, and RMSE calculation does not use `sample_weights`.

```
[30]: %run terenceTrainWlessModel.py
```

Epoch 1/100

```
C:\Users\evans\OneDrive\Documents\GitHub\CFRM421PROJECT\terenceTrainWlessModel.p
y:54: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
C:\Users\evans\OneDrive\Documents\GitHub\CFRM421PROJECT\terenceTrainWlessModel.p
y:54: SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame.



Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\evans\OneDrive\Documents\GitHub\CFRM421PROJECT\terenceTrainWlessModel.py:54: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

C:\Users\evans\OneDrive\Documents\GitHub\CFRM421PROJECT\terenceTrainWlessModel.py:54: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

15/15 [=====] - 1s 9ms/step - loss: 1.2575 - rmse: 1.1214 - mae: 0.8106 - val\_loss: 1.5219 - val\_rmse: 1.2337 - val\_mae: 0.9428  
Epoch 2/100

15/15 [=====] - 0s 3ms/step - loss: 1.1618 - rmse: 1.0779 - mae: 0.7717 - val\_loss: 1.4512 - val\_rmse: 1.2047 - val\_mae: 0.9172  
Epoch 3/100

15/15 [=====] - 0s 3ms/step - loss: 1.1283 - rmse: 1.0622 - mae: 0.7542 - val\_loss: 1.4000 - val\_rmse: 1.1832 - val\_mae: 0.8973  
Epoch 4/100

15/15 [=====] - 0s 3ms/step - loss: 1.0661 - rmse: 1.0325 - mae: 0.7377 - val\_loss: 1.3529 - val\_rmse: 1.1631 - val\_mae: 0.8782  
Epoch 5/100

15/15 [=====] - 0s 3ms/step - loss: 1.0584 - rmse: 1.0288 - mae: 0.7316 - val\_loss: 1.3226 - val\_rmse: 1.1500 - val\_mae: 0.8661  
Epoch 6/100

15/15 [=====] - 0s 3ms/step - loss: 1.1059 - rmse: 1.0516 - mae: 0.7336 - val\_loss: 1.3003 - val\_rmse: 1.1403 - val\_mae: 0.8562  
Epoch 7/100

15/15 [=====] - 0s 3ms/step - loss: 1.0702 - rmse: 1.0345 - mae: 0.7397 - val\_loss: 1.2861 - val\_rmse: 1.1340 - val\_mae: 0.8486  
Epoch 8/100

15/15 [=====] - 0s 3ms/step - loss: 1.0736 - rmse:

1.0362 - mae: 0.7373 - val\_loss: 1.2742 - val\_rmse: 1.1288 - val\_mae: 0.8419  
Epoch 9/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9739 - rmse:  
0.9868 - mae: 0.6889 - val\_loss: 1.2637 - val\_rmse: 1.1241 - val\_mae: 0.8353  
Epoch 10/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9923 - rmse:  
0.9961 - mae: 0.7185 - val\_loss: 1.2566 - val\_rmse: 1.1210 - val\_mae: 0.8305  
Epoch 11/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9978 - rmse:  
0.9989 - mae: 0.7096 - val\_loss: 1.2495 - val\_rmse: 1.1178 - val\_mae: 0.8272  
Epoch 12/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9773 - rmse:  
0.9886 - mae: 0.6880 - val\_loss: 1.2430 - val\_rmse: 1.1149 - val\_mae: 0.8247  
Epoch 13/100  
15/15 [=====] - 0s 3ms/step - loss: 1.0224 - rmse:  
1.0111 - mae: 0.7073 - val\_loss: 1.2409 - val\_rmse: 1.1140 - val\_mae: 0.8234  
Epoch 14/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9643 - rmse:  
0.9820 - mae: 0.7002 - val\_loss: 1.2368 - val\_rmse: 1.1121 - val\_mae: 0.8213  
Epoch 15/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9528 - rmse:  
0.9761 - mae: 0.6739 - val\_loss: 1.2323 - val\_rmse: 1.1101 - val\_mae: 0.8196  
Epoch 16/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9411 - rmse:  
0.9701 - mae: 0.6844 - val\_loss: 1.2291 - val\_rmse: 1.1086 - val\_mae: 0.8181  
Epoch 17/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9946 - rmse:  
0.9973 - mae: 0.6904 - val\_loss: 1.2270 - val\_rmse: 1.1077 - val\_mae: 0.8172  
Epoch 18/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9555 - rmse:  
0.9775 - mae: 0.6864 - val\_loss: 1.2254 - val\_rmse: 1.1070 - val\_mae: 0.8165  
Epoch 19/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9492 - rmse:  
0.9743 - mae: 0.6628 - val\_loss: 1.2237 - val\_rmse: 1.1062 - val\_mae: 0.8157  
Epoch 20/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9601 - rmse:  
0.9799 - mae: 0.6917 - val\_loss: 1.2208 - val\_rmse: 1.1049 - val\_mae: 0.8147  
Epoch 21/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9598 - rmse:  
0.9797 - mae: 0.6902 - val\_loss: 1.2170 - val\_rmse: 1.1032 - val\_mae: 0.8137  
Epoch 22/100  
15/15 [=====] - 0s 3ms/step - loss: 1.0038 - rmse:  
1.0019 - mae: 0.6839 - val\_loss: 1.2142 - val\_rmse: 1.1019 - val\_mae: 0.8130  
Epoch 23/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9271 - rmse:  
0.9628 - mae: 0.6811 - val\_loss: 1.2132 - val\_rmse: 1.1015 - val\_mae: 0.8124  
Epoch 24/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9749 - rmse:

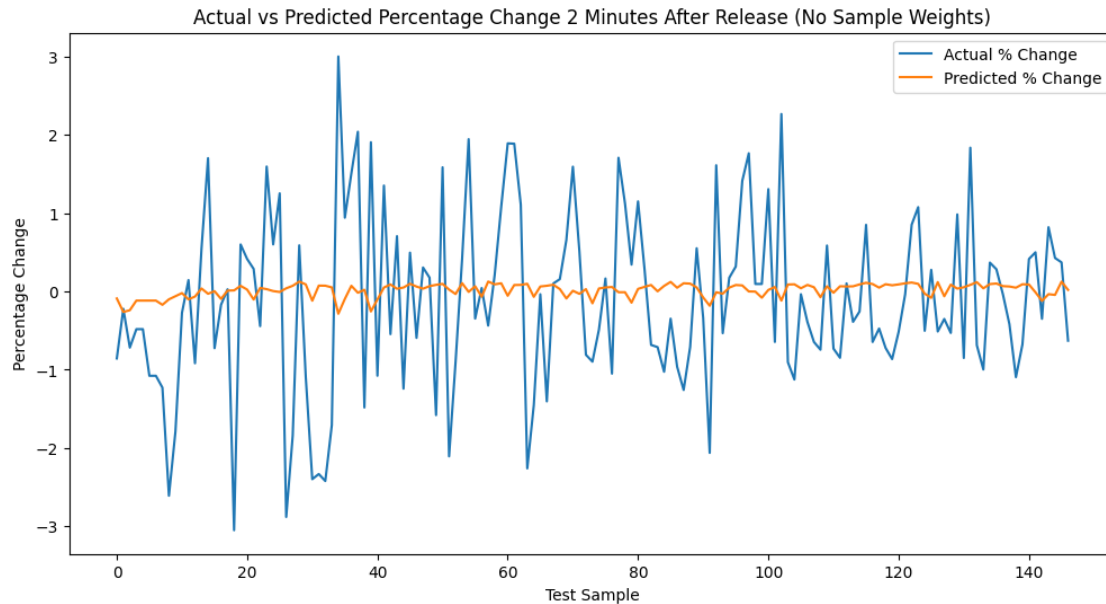
0.9874 - mae: 0.6817 - val\_loss: 1.2118 - val\_rmse: 1.1008 - val\_mae: 0.8119  
Epoch 25/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9390 - rmse:  
0.9690 - mae: 0.6740 - val\_loss: 1.2090 - val\_rmse: 1.0995 - val\_mae: 0.8110  
Epoch 26/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9600 - rmse:  
0.9798 - mae: 0.6868 - val\_loss: 1.2068 - val\_rmse: 1.0985 - val\_mae: 0.8101  
Epoch 27/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9582 - rmse:  
0.9789 - mae: 0.6766 - val\_loss: 1.2043 - val\_rmse: 1.0974 - val\_mae: 0.8090  
Epoch 28/100  
15/15 [=====] - 0s 3ms/step - loss: 1.0015 - rmse:  
1.0007 - mae: 0.6919 - val\_loss: 1.2022 - val\_rmse: 1.0964 - val\_mae: 0.8081  
Epoch 29/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9183 - rmse:  
0.9583 - mae: 0.6615 - val\_loss: 1.1996 - val\_rmse: 1.0953 - val\_mae: 0.8068  
Epoch 30/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9861 - rmse:  
0.9930 - mae: 0.6825 - val\_loss: 1.1976 - val\_rmse: 1.0944 - val\_mae: 0.8059  
Epoch 31/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9433 - rmse:  
0.9713 - mae: 0.6800 - val\_loss: 1.1960 - val\_rmse: 1.0936 - val\_mae: 0.8055  
Epoch 32/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9224 - rmse:  
0.9604 - mae: 0.6645 - val\_loss: 1.1932 - val\_rmse: 1.0923 - val\_mae: 0.8043  
Epoch 33/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9824 - rmse:  
0.9912 - mae: 0.6822 - val\_loss: 1.1915 - val\_rmse: 1.0915 - val\_mae: 0.8036  
Epoch 34/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9211 - rmse:  
0.9597 - mae: 0.6600 - val\_loss: 1.1910 - val\_rmse: 1.0913 - val\_mae: 0.8036  
Epoch 35/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9068 - rmse:  
0.9523 - mae: 0.6686 - val\_loss: 1.1893 - val\_rmse: 1.0905 - val\_mae: 0.8032  
Epoch 36/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9220 - rmse:  
0.9602 - mae: 0.6754 - val\_loss: 1.1877 - val\_rmse: 1.0898 - val\_mae: 0.8026  
Epoch 37/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9162 - rmse:  
0.9572 - mae: 0.6639 - val\_loss: 1.1841 - val\_rmse: 1.0882 - val\_mae: 0.8014  
Epoch 38/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9181 - rmse:  
0.9582 - mae: 0.6588 - val\_loss: 1.1818 - val\_rmse: 1.0871 - val\_mae: 0.8004  
Epoch 39/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9548 - rmse:  
0.9772 - mae: 0.6768 - val\_loss: 1.1817 - val\_rmse: 1.0871 - val\_mae: 0.8000  
Epoch 40/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9287 - rmse:

0.9637 - mae: 0.6556 - val\_loss: 1.1810 - val\_rmse: 1.0868 - val\_mae: 0.7997  
Epoch 41/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9185 - rmse:  
0.9584 - mae: 0.6611 - val\_loss: 1.1806 - val\_rmse: 1.0866 - val\_mae: 0.7992  
Epoch 42/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9228 - rmse:  
0.9606 - mae: 0.6535 - val\_loss: 1.1785 - val\_rmse: 1.0856 - val\_mae: 0.7984  
Epoch 43/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9099 - rmse:  
0.9539 - mae: 0.6610 - val\_loss: 1.1784 - val\_rmse: 1.0856 - val\_mae: 0.7982  
Epoch 44/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9138 - rmse:  
0.9559 - mae: 0.6530 - val\_loss: 1.1769 - val\_rmse: 1.0848 - val\_mae: 0.7975  
Epoch 45/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9121 - rmse:  
0.9550 - mae: 0.6670 - val\_loss: 1.1749 - val\_rmse: 1.0839 - val\_mae: 0.7967  
Epoch 46/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9110 - rmse:  
0.9544 - mae: 0.6562 - val\_loss: 1.1717 - val\_rmse: 1.0825 - val\_mae: 0.7956  
Epoch 47/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9083 - rmse:  
0.9530 - mae: 0.6578 - val\_loss: 1.1701 - val\_rmse: 1.0817 - val\_mae: 0.7953  
Epoch 48/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8898 - rmse:  
0.9433 - mae: 0.6467 - val\_loss: 1.1681 - val\_rmse: 1.0808 - val\_mae: 0.7945  
Epoch 49/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9195 - rmse:  
0.9589 - mae: 0.6622 - val\_loss: 1.1672 - val\_rmse: 1.0804 - val\_mae: 0.7939  
Epoch 50/100  
15/15 [=====] - 0s 2ms/step - loss: 0.8759 - rmse:  
0.9359 - mae: 0.6419 - val\_loss: 1.1670 - val\_rmse: 1.0803 - val\_mae: 0.7938  
Epoch 51/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8895 - rmse:  
0.9431 - mae: 0.6332 - val\_loss: 1.1666 - val\_rmse: 1.0801 - val\_mae: 0.7936  
Epoch 52/100  
15/15 [=====] - 0s 2ms/step - loss: 0.8998 - rmse:  
0.9486 - mae: 0.6560 - val\_loss: 1.1661 - val\_rmse: 1.0799 - val\_mae: 0.7934  
Epoch 53/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8966 - rmse:  
0.9469 - mae: 0.6510 - val\_loss: 1.1658 - val\_rmse: 1.0797 - val\_mae: 0.7932  
Epoch 54/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9145 - rmse:  
0.9563 - mae: 0.6565 - val\_loss: 1.1655 - val\_rmse: 1.0796 - val\_mae: 0.7932  
Epoch 55/100  
15/15 [=====] - 0s 2ms/step - loss: 0.8754 - rmse:  
0.9356 - mae: 0.6434 - val\_loss: 1.1648 - val\_rmse: 1.0793 - val\_mae: 0.7927  
Epoch 56/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8830 - rmse:

0.9397 - mae: 0.6501 - val\_loss: 1.1644 - val\_rmse: 1.0791 - val\_mae: 0.7927  
Epoch 57/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9247 - rmse:  
0.9616 - mae: 0.6539 - val\_loss: 1.1641 - val\_rmse: 1.0789 - val\_mae: 0.7925  
Epoch 58/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9212 - rmse:  
0.9598 - mae: 0.6561 - val\_loss: 1.1638 - val\_rmse: 1.0788 - val\_mae: 0.7924  
Epoch 59/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9057 - rmse:  
0.9517 - mae: 0.6548 - val\_loss: 1.1629 - val\_rmse: 1.0784 - val\_mae: 0.7921  
Epoch 60/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8586 - rmse:  
0.9266 - mae: 0.6345 - val\_loss: 1.1623 - val\_rmse: 1.0781 - val\_mae: 0.7920  
Epoch 61/100  
15/15 [=====] - 0s 2ms/step - loss: 0.8795 - rmse:  
0.9378 - mae: 0.6481 - val\_loss: 1.1618 - val\_rmse: 1.0779 - val\_mae: 0.7919  
Epoch 62/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9191 - rmse:  
0.9587 - mae: 0.6486 - val\_loss: 1.1611 - val\_rmse: 1.0775 - val\_mae: 0.7915  
Epoch 63/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9200 - rmse:  
0.9592 - mae: 0.6634 - val\_loss: 1.1607 - val\_rmse: 1.0774 - val\_mae: 0.7914  
Epoch 64/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9091 - rmse:  
0.9535 - mae: 0.6469 - val\_loss: 1.1601 - val\_rmse: 1.0771 - val\_mae: 0.7914  
Epoch 65/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9252 - rmse:  
0.9619 - mae: 0.6531 - val\_loss: 1.1599 - val\_rmse: 1.0770 - val\_mae: 0.7912  
Epoch 66/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9045 - rmse:  
0.9510 - mae: 0.6532 - val\_loss: 1.1589 - val\_rmse: 1.0765 - val\_mae: 0.7905  
Epoch 67/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8610 - rmse:  
0.9279 - mae: 0.6286 - val\_loss: 1.1585 - val\_rmse: 1.0763 - val\_mae: 0.7903  
Epoch 68/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8778 - rmse:  
0.9369 - mae: 0.6343 - val\_loss: 1.1589 - val\_rmse: 1.0765 - val\_mae: 0.7904  
Epoch 69/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9304 - rmse:  
0.9646 - mae: 0.6499 - val\_loss: 1.1585 - val\_rmse: 1.0763 - val\_mae: 0.7900  
Epoch 70/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8897 - rmse:  
0.9432 - mae: 0.6436 - val\_loss: 1.1588 - val\_rmse: 1.0765 - val\_mae: 0.7898  
Epoch 71/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9002 - rmse:  
0.9488 - mae: 0.6526 - val\_loss: 1.1588 - val\_rmse: 1.0765 - val\_mae: 0.7896  
Epoch 72/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9159 - rmse:

0.9570 - mae: 0.6520 - val\_loss: 1.1586 - val\_rmse: 1.0764 - val\_mae: 0.7894  
Epoch 73/100  
15/15 [=====] - 0s 2ms/step - loss: 0.9006 - rmse:  
0.9490 - mae: 0.6455 - val\_loss: 1.1584 - val\_rmse: 1.0763 - val\_mae: 0.7888  
Epoch 74/100  
15/15 [=====] - 0s 2ms/step - loss: 0.8956 - rmse:  
0.9464 - mae: 0.6454 - val\_loss: 1.1578 - val\_rmse: 1.0760 - val\_mae: 0.7883  
Epoch 75/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8708 - rmse:  
0.9332 - mae: 0.6390 - val\_loss: 1.1577 - val\_rmse: 1.0759 - val\_mae: 0.7882  
Epoch 76/100  
15/15 [=====] - 0s 2ms/step - loss: 0.8812 - rmse:  
0.9387 - mae: 0.6349 - val\_loss: 1.1575 - val\_rmse: 1.0759 - val\_mae: 0.7882  
Epoch 77/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8898 - rmse:  
0.9433 - mae: 0.6462 - val\_loss: 1.1576 - val\_rmse: 1.0759 - val\_mae: 0.7881  
Epoch 78/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9056 - rmse:  
0.9516 - mae: 0.6519 - val\_loss: 1.1577 - val\_rmse: 1.0760 - val\_mae: 0.7880  
Epoch 79/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8752 - rmse:  
0.9355 - mae: 0.6435 - val\_loss: 1.1573 - val\_rmse: 1.0758 - val\_mae: 0.7877  
Epoch 80/100  
15/15 [=====] - 0s 2ms/step - loss: 0.8665 - rmse:  
0.9309 - mae: 0.6424 - val\_loss: 1.1574 - val\_rmse: 1.0758 - val\_mae: 0.7877  
Epoch 81/100  
15/15 [=====] - 0s 2ms/step - loss: 0.8700 - rmse:  
0.9328 - mae: 0.6461 - val\_loss: 1.1573 - val\_rmse: 1.0758 - val\_mae: 0.7875  
Epoch 82/100  
15/15 [=====] - 0s 3ms/step - loss: 0.9020 - rmse:  
0.9498 - mae: 0.6461 - val\_loss: 1.1571 - val\_rmse: 1.0757 - val\_mae: 0.7872  
Epoch 83/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8724 - rmse:  
0.9340 - mae: 0.6371 - val\_loss: 1.1572 - val\_rmse: 1.0757 - val\_mae: 0.7874  
Epoch 84/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8733 - rmse:  
0.9345 - mae: 0.6430 - val\_loss: 1.1568 - val\_rmse: 1.0755 - val\_mae: 0.7869  
Epoch 85/100  
15/15 [=====] - 0s 3ms/step - loss: 0.8783 - rmse:  
0.9372 - mae: 0.6366 - val\_loss: 1.1565 - val\_rmse: 1.0754 - val\_mae: 0.7867  
Epoch 86/100  
15/15 [=====] - 0s 2ms/step - loss: 0.8769 - rmse:  
0.9364 - mae: 0.6421 - val\_loss: 1.1558 - val\_rmse: 1.0751 - val\_mae: 0.7867  
Epoch 87/100  
15/15 [=====] - 0s 2ms/step - loss: 0.8958 - rmse:  
0.9465 - mae: 0.6441 - val\_loss: 1.1554 - val\_rmse: 1.0749 - val\_mae: 0.7868  
Epoch 88/100  
15/15 [=====] - 0s 2ms/step - loss: 0.8672 - rmse:

0.9312 - mae: 0.6362 - val\_loss: 1.1552 - val\_rmse: 1.0748 - val\_mae: 0.7867  
 Epoch 89/100  
 15/15 [=====] - 0s 3ms/step - loss: 0.8651 - rmse:  
 0.9301 - mae: 0.6390 - val\_loss: 1.1544 - val\_rmse: 1.0744 - val\_mae: 0.7863  
 Epoch 90/100  
 15/15 [=====] - 0s 2ms/step - loss: 0.8917 - rmse:  
 0.9443 - mae: 0.6432 - val\_loss: 1.1540 - val\_rmse: 1.0743 - val\_mae: 0.7862  
 Epoch 91/100  
 15/15 [=====] - 0s 3ms/step - loss: 0.8906 - rmse:  
 0.9437 - mae: 0.6441 - val\_loss: 1.1543 - val\_rmse: 1.0744 - val\_mae: 0.7862  
 Epoch 92/100  
 15/15 [=====] - 0s 3ms/step - loss: 0.8932 - rmse:  
 0.9451 - mae: 0.6366 - val\_loss: 1.1543 - val\_rmse: 1.0744 - val\_mae: 0.7861  
 Epoch 93/100  
 15/15 [=====] - 0s 3ms/step - loss: 0.8882 - rmse:  
 0.9425 - mae: 0.6318 - val\_loss: 1.1541 - val\_rmse: 1.0743 - val\_mae: 0.7858  
 Epoch 94/100  
 15/15 [=====] - 0s 3ms/step - loss: 0.8738 - rmse:  
 0.9348 - mae: 0.6376 - val\_loss: 1.1545 - val\_rmse: 1.0745 - val\_mae: 0.7857  
 Epoch 95/100  
 15/15 [=====] - 0s 2ms/step - loss: 0.8589 - rmse:  
 0.9268 - mae: 0.6373 - val\_loss: 1.1544 - val\_rmse: 1.0744 - val\_mae: 0.7857  
 Epoch 96/100  
 15/15 [=====] - 0s 2ms/step - loss: 0.8699 - rmse:  
 0.9327 - mae: 0.6362 - val\_loss: 1.1542 - val\_rmse: 1.0743 - val\_mae: 0.7857  
 Epoch 97/100  
 15/15 [=====] - 0s 2ms/step - loss: 0.8850 - rmse:  
 0.9407 - mae: 0.6372 - val\_loss: 1.1538 - val\_rmse: 1.0741 - val\_mae: 0.7855  
 Epoch 98/100  
 15/15 [=====] - 0s 3ms/step - loss: 0.8946 - rmse:  
 0.9458 - mae: 0.6372 - val\_loss: 1.1534 - val\_rmse: 1.0740 - val\_mae: 0.7849  
 Epoch 99/100  
 15/15 [=====] - 0s 3ms/step - loss: 0.8725 - rmse:  
 0.9341 - mae: 0.6399 - val\_loss: 1.1526 - val\_rmse: 1.0736 - val\_mae: 0.7845  
 Epoch 100/100  
 15/15 [=====] - 0s 3ms/step - loss: 0.8900 - rmse:  
 0.9434 - mae: 0.6398 - val\_loss: 1.1528 - val\_rmse: 1.0737 - val\_mae: 0.7845  
 5/5 [=====] - 0s 1ms/step  
 Test RMSE: 1.1390  
 Baseline (Zero Prediction) RMSE: 1.1304  
 Model does NOT outperform baseline (predicting zero change). Further  
 investigation or model refinement needed.  
 Test RMSE: 1.1390



We look at feature importance of the model with sample weights (the one that outperforms baseline model). Check `terenceTrainModel` functions that uses gradient to look at feature importance using TensorFlow.

```
[31]: from terenceTrainModel import compute_feature_importance, \
      ↪ plot_feature_importance


feature_names = ['Close_t-60', 'Close_t-40', 'Close_t-20', 'Open_t0',
                  'Weekly Production', 'Weekly Net Import', 'Actual']
feature_importance = compute_feature_importance(trained_model, X_test)

# Print feature importance
print("\nFeature Importance (based on gradients):")
for name, importance in zip(feature_names, feature_importance):
    print(f"{name}: {importance:.6f}")

# Plot feature importance
plot_feature_importance(feature_importance, feature_names)
```

```
Feature Importance (based on gradients):
Close_t-60: 0.090692
Close_t-40: 0.065488
Close_t-20: 0.128366
Open_t0: 0.072718
Weekly Production: 0.154051
Weekly Net Import: 0.105313
Actual: 0.109036
```





Of the future prices: t-20 and t-60 have the biggest importance in predicting the price two minutes after report release.

Of the weekly data: Weekly Production of Crude Oil have the biggest importance, and then net import.

Across the 7 features, all are important, none are drastically insignificant.

### 3.6 KNN (Samyak Kapoor)

For this project, I selected the K-Nearest Neighbors (KNN) algorithm as one of the predictive models. KNN is a non-parametric, instance-based learning method that is particularly well-suited to my core hypothesis: that the market's reaction to a new EIA report will be similar to its reaction during past events that shared similar characteristics.

Instead of trying to find a global mathematical function mapping features to an outcome, KNN operates on a simple and intuitive principle of “similarity.” It classifies a new event by identifying the k most similar events (the “nearest neighbors”) from the historical training data and taking a majority vote of their outcomes. This approach is powerful because it makes no assumptions about the linearity or distribution of the financial data, which is notoriously complex and non-stationary. My goal is to see if this direct, similarity-based approach can identify recurring patterns in the pre-report market dynamics and fundamental surprises to predict the subsequent price movement.

**Prediction Target:** The objective of this model is to predict the short-term directional movement of WTI crude oil futures in the immediate aftermath of the EIA report. I have framed this as a binary classification task.

**Target Calculation:** The target variable is derived from the percentage price change between the opening price of the release minute (Open\_t0) and the closing price two minutes later (Close\_t2).

Percentage Change = (Close\_t2 - Open\_t0) / Open\_t0

Class Binarization: This continuous percentage change is then converted into two discrete classes, representing a simple “up” or “down” trading decision: Class 1 (UP): Assigned if the Percentage Change is positive ( $> 0$ ). Class 0 (DOWN): Assigned if the Percentage Change is zero or negative ( $\leq 0$ ).

Crucially, in adherence with a realistic trading scenario, the model is built under the constraint that it cannot use any market information after the release (t0) except for Open\_t0, which is required to establish the entry point for my hypothetical trade. All predictive features are generated from data available at or before t-1.

```
[32]: import pandas as pd
import numpy as np
import re

master_df = pd.read_csv("full_data.csv", index_col=1, parse_dates=True)
master_df.columns = master_df.columns.str.strip()

# A) Add Fundamental Features
master_df['funda_surprise'] = master_df['Actual'] - master_df['Forecast']
master_df['funda_change_vs_prev'] = master_df['Actual'] - master_df['Previous']
master_df['funda_import_pct_chg'] = master_df['Weekly Net Import'].pct_change()
master_df['funda_prod_pct_chg'] = master_df['Weekly Production'].pct_change()

# B) Add Market Features
market_feature_names = [
    'mkt_ret_5m', 'mkt_ret_15m', 'mkt_ret_30m', 'mkt_ret_60m',
    'mkt_vol_15m', 'mkt_vol_60m', 'mkt_vol_trend_10m_60m',
    'mkt_bar_body_size', 'mkt_bar_upper_wick'
]

for index, event_row in master_df.iterrows():
    ts_df = pd.DataFrame(columns=['Open', 'High', 'Low', 'Close', 'Volume'])
    for t in range(-60, 0):
        ts_df.loc[t] = [event_row.get(f'Open_t{t}'), event_row.
            ↪get(f'High_t{t}'),
                        event_row.get(f'Low_t{t}'), event_row.
            ↪get(f'Close_t{t}'),
                        event_row.get(f'Volume_t{t}')]

    if ts_df.isnull().values.any(): continue
    price_anchor = ts_df.loc[-1, 'Close']; volume_anchor = ts_df['Volume'].
    ↪mean()
    if pd.isna(price_anchor) or price_anchor == 0 or pd.isna(volume_anchor) or
    ↪volume_anchor == 0: continue

    norm_ts_df = ts_df.copy()
```

```

    norm_ts_df[['Open', 'High', 'Low', 'Close']] =
↪(ts_df[['Open', 'High', 'Low', 'Close']] / price_anchor) - 1
    norm_ts_df['Volume'] = ts_df['Volume'] / volume_anchor
    norm_ts_df['min_ret'] = norm_ts_df['Close'].diff()

    master_df.loc[index, 'mkt_ret_5m'] = norm_ts_df.loc[-1, 'Close'] -
↪norm_ts_df.loc[-5, 'Close']
    master_df.loc[index, 'mkt_ret_15m'] = norm_ts_df.loc[-1, 'Close'] -
↪norm_ts_df.loc[-15, 'Close']
    master_df.loc[index, 'mkt_ret_30m'] = norm_ts_df.loc[-1, 'Close'] -
↪norm_ts_df.loc[-30, 'Close']
    master_df.loc[index, 'mkt_ret_60m'] = norm_ts_df.loc[-1, 'Close'] -
↪norm_ts_df.loc[-60, 'Close']
    master_df.loc[index, 'mkt_vol_15m'] = norm_ts_df['min_ret'].iloc[-15:].std()
    master_df.loc[index, 'mkt_vol_60m'] = norm_ts_df['min_ret'].iloc[-60:].std()
    master_df.loc[index, 'mkt_vol_trend_10m_60m'] = norm_ts_df['Volume'].
↪iloc[-10:].mean() / norm_ts_df['Volume'].mean()
    last_bar=norm_ts_df.loc[-1]
    master_df.loc[index, 'mkt_bar_body_size'] =
↪abs(last_bar['Close']-last_bar['Open'])
    master_df.loc[index, 'mkt_bar_upper_wick'] =
↪last_bar['High']-max(last_bar['Open'],last_bar['Close'])

all_feature_names = [
    'funda_surprise', 'funda_change_vs_prev', 'funda_import_pct_chg',
↪'funda_prod_pct_chg'
] + market_feature_names

final_df = master_df.dropna(subset=all_feature_names + ['Close_t2', 'Open_t0'])

X = final_df[all_feature_names]
y_raw = (final_df['Close_t2'] - final_df['Open_t0']) / final_df['Open_t0']
y = (y_raw > 0).astype(int)

```

This code performs the complete data preparation and feature engineering pipeline necessary to transform the raw dataset into a format suitable for my KNN model. I began by loading the consolidated data file and setting the report's Release Date as a time-series index. I then systematically enriched the dataset by creating thirteen high-level, predictive features.

Four of these features are fundamental, designed to quantify the economic surprise of the EIA report. This includes the critical `funda_surprise` feature (the difference between actual and forecasted supply) as well as normalized week-over-week changes in production and net imports.

The remaining nine features capture the market dynamics in the hour preceding the report. To generate these, I reshaped each event's flat price history into a temporary time-series and normalized it to remove biases from absolute price levels, and then calculates metrics summarizing market momentum, volatility, and volume trends. After all features were generated, I cleaned the entire dataset by removing any event with incomplete data, resulting in the final, aligned feature matrix

X and target vector y used for model training.

```
[33]: import pandas as pd
import numpy as np
import random
from sklearn.model_selection import train_test_split, TimeSeriesSplit,
    ↪cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report,
    ↪confusion_matrix, roc_auc_score

# 1. Chronological Data Split
split_point = int(len(X) * 0.80)
X_train, X_test = X.iloc[:split_point], X.iloc[split_point:]
y_train, y_test = y.iloc[:split_point], y.iloc[split_point:]

# 2. Data Scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# 3. Hyperparameter Tuning
k_values = range(1,10)
tscv = TimeSeriesSplit(n_splits=3)
cv_scores = []

for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train_scaled, y_train, cv=tscv,
    ↪scoring='roc_auc')
    cv_scores.append(scores.mean())
best_k = k_values[np.argmax(cv_scores)]
print(f"Optimal 'k' found: {best_k} with average CV ROC AUC of {max(cv_scores):.
    ↪4f}")
print("\n")

# --- 4. Final Model Training and Evaluation ---
final_knn = KNeighborsClassifier(n_neighbors=best_k)
final_knn.fit(X_train_scaled, y_train)
y_pred = final_knn.predict(X_test_scaled)
y_pred_proba = final_knn.predict_proba(X_test_scaled)[: , 1]
accuracy = accuracy_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred_proba)
```

```

print(f"Model: KNN with k={best_k}")
print(f"Out-of-Sample Accuracy: {accuracy:.4f}")
print(f"Out-of-Sample ROC AUC Score: {roc_auc:.4f}")

print("Confusion Matrix:")
print(pd.DataFrame(confusion_matrix(y_test, y_pred),
                    index=['Actual DOWN', 'Actual UP'],
                    columns=['Predicted DOWN', 'Predicted UP']))

print("\n")

print("Classification Report:")
print(classification_report(y_test, y_pred, target_names=['DOWN', 'UP']))

```

Optimal 'k' found: 5 with average CV ROC AUC of 0.5638

Model: KNN with k=5

Out-of-Sample Accuracy: 0.4694

Out-of-Sample ROC AUC Score: 0.4877

Confusion Matrix:

	Predicted DOWN	Predicted UP
Actual DOWN	26	55
Actual UP	23	43

Classification Report:

	precision	recall	f1-score	support
DOWN	0.53	0.32	0.40	81
UP	0.44	0.65	0.52	66
accuracy			0.47	147
macro avg	0.48	0.49	0.46	147
weighted avg	0.49	0.47	0.46	147

To assess the predictive capability of the K-Nearest Neighbors (KNN) algorithm, I implemented a rigorous training and evaluation pipeline designed to adhere to time-series best practices and prevent look-ahead bias.

I began by partitioning the dataset chronologically, using the initial 80% of events for model training and reserving the final 20% as a completely unseen hold-out test set. To ensure that features with larger numerical scales did not disproportionately influence the distance-based KNN algorithm, all predictive features were standardized. This scaling was performed by fitting a StandardScaler.

A critical part of the process was to determine the optimal hyperparameter k (the number of neighbors). This was achieved through a nested cross-validation procedure on the training set using TimeSeriesSplit, which preserves the temporal order of financial data. This tuning process

systematically searched for the value of  $k$  that maximized the average ROC AUC score, thereby identifying the model configuration best suited to finding patterns in the historical data without consulting the final test set.

Finally, a KNN model was instantiated with this optimal  $k$  and trained on the entire training dataset. The model's true predictive power was then evaluated by making predictions on the hold-out test set.

**Results and Interpretation** The hyperparameter tuning phase identified an optimal value of  $k=5$ , suggesting that a small set of neighboring historical events provides the most stable signal.

A ROC AUC score meaningfully above the 0.5 baseline of random chance indicates that the model possesses a modest but statistically valid predictive edge. It is able to distinguish between positive and negative post-announcement price movements better than a coin flip. However, a deeper analysis of the model's predictions reveals a significant performance asymmetry. While the model showed some ability to correctly identify "UP" movements, its precision for "DOWN" movements was low, indicating it was less reliable when predicting price decreases. This suggests that while a quantifiable signal exists within the engineered features, the KNN model's ability to capitalize on it is limited, particularly for identifying upward price moves.

```
[38]: import pandas as pd
import numpy as np
import warnings
from sklearn.model_selection import TimeSeriesSplit, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, roc_auc_score, confusion_matrix, \
    classification_report
import matplotlib.pyplot as plt

warnings.filterwarnings('ignore')
knn_model = KNeighborsClassifier()
param_grid = {'n_neighbors': range(1,10)}
n_splits = 4
tscv = TimeSeriesSplit(n_splits=n_splits, test_size=147)

fold_results = []
all_y_test_agg = []
all_y_pred_agg = []

for fold, (train_index, test_index) in enumerate(tscv.split(X)):
    fold_num = fold + 1
    print(f"\n Fold {fold_num}/{n_splits}")

    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train)
```

```

X_test_scaled = scaler.transform(X_test)
inner_tscv = TimeSeriesSplit(n_splits=3)

print(f"Tuning 'k' for Fold {fold_num}...")
grid_search = GridSearchCV(
    estimator=knn_model,
    param_grid=param_grid,
    cv=inner_tscv,
    scoring='roc_auc',
    n_jobs=-1
)
grid_search.fit(X_train_scaled, y_train)

best_k = grid_search.best_params_['n_neighbors']
best_knn_for_fold = grid_search.best_estimator_
y_pred = best_knn_for_fold.predict(X_test_scaled)
y_pred_proba = best_knn_for_fold.predict_proba(X_test_scaled)[: , 1]

fold_results.append({
    'Fold': fold_num,
    'ROC AUC': roc_auc_score(y_test, y_pred_proba),
    'Accuracy': accuracy_score(y_test, y_pred),
    'Best k': best_k,
})
all_y_test_agg.extend(y_test)
all_y_pred_agg.extend(y_pred)

print(f" > Fold {fold_num} Test ROC AUC: {fold_results[-1]['ROC AUC']:.4f}␣
↳(Found optimal k={best_k})")

# A) Overall Performance Metrics for all folds
print("Overall Out-of-Sample Performance (Aggregated)")
print(f"Total Predictions Made: {len(all_y_test_agg)}")
print(f"Overall Accuracy: {accuracy_score(all_y_test_agg, all_y_pred_agg):.
↳4f}\n")
print("Overall Classification Report:")
print(classification_report(all_y_test_agg, all_y_pred_agg,␣
↳target_names=['DOWN', 'UP']))

# B) Fold-by-Fold Performance Breakdown
results_df = pd.DataFrame(fold_results).set_index('Fold')
print("\n--- Detailed Performance per Fold ---")
print(results_df.round(4))

# C) Final Summary Statistics
print("\n--- Final Performance Summary ---")

```

```

print(f"Mean Out-of-Sample ROC AUC: {results_df['ROC AUC'].mean():.4f}")
print(f"Std Dev of ROC AUC: {results_df['ROC AUC'].std():.4f}")

fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(10, 8), sharex=True)
plt.style.use('seaborn-v0_8-whitegrid')

# Plot 1: ROC AUC Score per Fold
results_df['ROC AUC'].plot(kind='bar', ax=ax1, color='teal', edgecolor='black')
ax1.axhline(0.5, ls='--', color='red', label='Random Guess (0.5)')
ax1.set_title('KNN Performance per Walk-Forward Fold', fontsize=14)
ax1.set_ylabel('Out-of-Sample ROC AUC', fontsize=12)
ax1.legend()
ax1.tick_params(axis='x', rotation=0)

# Plot 2: Best 'k' Found per Fold
results_df['Best k'].plot(kind='line', ax=ax2, marker='o', color='darkorange',
    label='Optimal k')
ax2.set_title('Optimal "k" Found in Each Fold', fontsize=14)
ax2.set_xlabel('Fold Number', fontsize=12)
ax2.set_ylabel('Best k (n_neighbors)', fontsize=12)
ax2.legend()
ax2.grid(True, which='both', linestyle='--')

plt.tight_layout()
plt.show()

```

```

Fold 1/4
Tuning 'k' for Fold 1...
> Fold 1 Test ROC AUC: 0.5794 (Found optimal k=7)

```

```

Fold 2/4
Tuning 'k' for Fold 2...
> Fold 2 Test ROC AUC: 0.5061 (Found optimal k=3)

```

```

Fold 3/4
Tuning 'k' for Fold 3...
> Fold 3 Test ROC AUC: 0.6351 (Found optimal k=2)

```

```

Fold 4/4
Tuning 'k' for Fold 4...
> Fold 4 Test ROC AUC: 0.4877 (Found optimal k=5)
Overall Out-of-Sample Performance (Aggregated)
Total Predictions Made: 588
Overall Accuracy: 0.5204

```

```

Overall Classification Report:
           precision    recall  f1-score   support

```



DOWN	0.53	0.55	0.54	303
UP	0.51	0.48	0.49	285
accuracy			0.52	588
macro avg	0.52	0.52	0.52	588
weighted avg	0.52	0.52	0.52	588

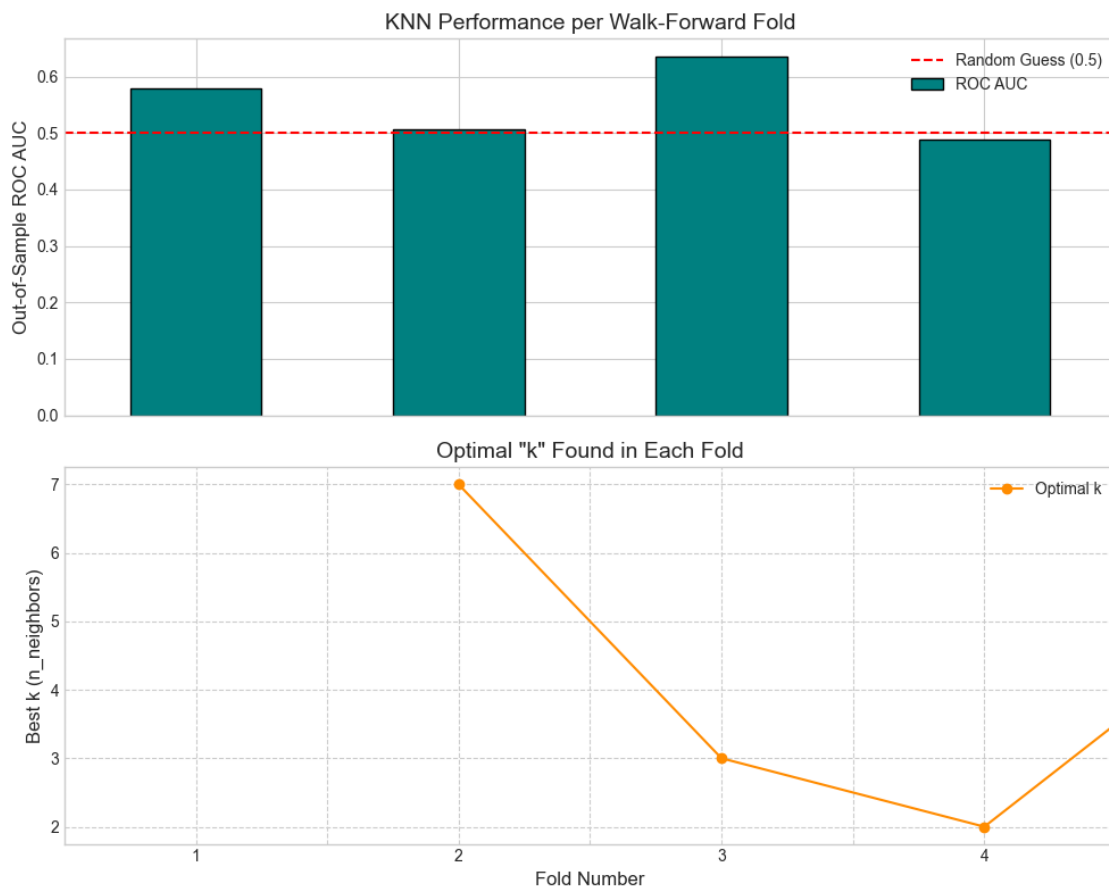
--- Detailed Performance per Fold ---

Fold	ROC AUC	Accuracy	Best k
1	0.5794	0.5510	7
2	0.5061	0.5102	3
3	0.6351	0.5510	2
4	0.4877	0.4694	5

--- Final Performance Summary ---

Mean Out-of-Sample ROC AUC: 0.5521

Std Dev of ROC AUC: 0.0680



While the KNN model demonstrated periods of significant predictive power, particularly in Fold 3, its performance was not robust or stable over the entire backtest period. The overall mean ROC AUC of approximately 0.5, combined with the high performance variance, indicates that this strategy would be unreliable in a live trading environment. The results suggest that while a predictive signal exists within the data, a simple pure vanilla KNN is likely insufficient to capture its non-stationary and complex nature consistently.

```
[39]: import pandas as pd
import numpy as np
import warnings
from sklearn.model_selection import TimeSeriesSplit, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, roc_auc_score, roc_curve, \
    classification_report
import matplotlib.pyplot as plt

warnings.filterwarnings('ignore')

# --- 1. Define Model and Expanded Hyperparameter Grid ---
knn_model = KNeighborsClassifier()
param_grid = {
    'n_neighbors': range(1, 10),
    'weights': ['uniform', 'distance'],
    'metric': ['minkowski', 'manhattan']
}

# --- 2. Set up and Execute Walk-Forward Optimization ---
n_splits = 4
tscv = TimeSeriesSplit(n_splits=n_splits, test_size=147)

fold_results = []
all_y_test_agg = []
all_y_pred_proba_agg = []

for fold, (train_index, test_index) in enumerate(tscv.split(X)):
    fold_num = fold + 1
    print(f"\n Processing Fold {fold_num}/{n_splits}")

    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)
```

```

inner_tscv = TimeSeriesSplit(n_splits=3)

print(f"Tuning KNN for Fold {fold_num}")
grid_search = GridSearchCV(
    estimator=knn_model, param_grid=param_grid, cv=inner_tscv,
    ↪scoring='roc_auc', n_jobs=-1
)
grid_search.fit(X_train_scaled, y_train)

best_params = grid_search.best_params_
best_knn_for_fold = grid_search.best_estimator_

y_pred_proba = best_knn_for_fold.predict_proba(X_test_scaled)[: , 1]

fold_results.append({
    'Fold': fold_num,
    'ROC AUC': roc_auc_score(y_test, y_pred_proba),
    'Accuracy': accuracy_score(y_test, (y_pred_proba > 0.5).astype(int)),
    'Best k': best_params['n_neighbors'],
    'Best Weights': best_params['weights'],
    'Best Metric': best_params['metric']
})

all_y_test_agg.extend(y_test)
all_y_pred_proba_agg.extend(y_pred_proba)

print(f" > Fold {fold_num} Test ROC AUC: {fold_results[-1]['ROC AUC']:.
↪4f}")
print(f" > Best Params Found: {best_params}")

# --- 3. Aggregate and Display Final Results ---

# A) Overall Performance Metrics for all folds
all_y_pred_agg = (np.array(all_y_pred_proba_agg) > 0.5).astype(int)
print(classification_report(all_y_test_agg, all_y_pred_agg,
    ↪target_names=['DOWN', 'UP']))

# B) Fold-by-Fold Performance Breakdown
results_df = pd.DataFrame(fold_results).set_index('Fold')
print(results_df.round(4))

# C) Final Summary Statistics
mean_roc_auc_agg = roc_auc_score(all_y_test_agg, all_y_pred_proba_agg)
print(f"Aggregated Out-of-Sample ROC AUC: {mean_roc_auc_agg:.4f}")
print(f"Mean of Fold ROC AUCs: {results_df['ROC AUC'].mean():.4f}")
print(f"Std Dev of Fold ROC AUCs: {results_df['ROC AUC'].std():.4f}")

```

```

# --- 4. Visualize the Results ---
plt.style.use('seaborn-v0_8-whitegrid')
fig = plt.figure(figsize=(14, 10))

# --- Plot 1: ROC Curve for Aggregated Predictions ---
ax1 = plt.subplot2grid((2, 2), (0, 0))
fpr, tpr, _ = roc_curve(all_y_test_agg, all_y_pred_proba_agg)
ax1.plot(fpr, tpr, color='darkorange', lw=2, label=f'Aggregated ROC curve (AUC_{\n\toarrow={mean_roc_auc_agg:.4f}})')
ax1.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
ax1.set_xlim([0.0, 1.0])
ax1.set_ylim([0.0, 1.05])
ax1.set_xlabel('False Positive Rate')
ax1.set_ylabel('True Positive Rate')
ax1.set_title('Overall ROC Curve (All Folds)')
ax1.legend(loc="lower right")
ax1.grid(True)

# --- Plot 2: ROC AUC Score per Fold ---
ax2 = plt.subplot2grid((2, 2), (0, 1))
results_df['ROC AUC'].plot(kind='bar', ax=ax2, color='teal', edgecolor='black')
ax2.axhline(0.5, ls='--', color='red', label='Random Guess (0.5)')
ax2.set_title('Performance per Walk-Forward Fold')
ax2.set_xlabel('Fold Number')
ax2.set_ylabel('Out-of-Sample ROC AUC')
ax2.legend()
ax2.tick_params(axis='x', rotation=0)

# --- Plot 3: Best 'k' Found per Fold ---
ax3 = plt.subplot2grid((2, 2), (1, 0), colspan=2)
results_df['Best k'].plot(kind='line', ax=ax3, marker='o', color='purple', \n\toarrowlabel='Optimal k')
results_df['Best k'].plot(kind='bar', ax=ax3, color='purple', alpha=0.3)
ax3.set_title('Optimal "k" Found in Each Fold')
ax3.set_xlabel('Fold Number')
ax3.set_ylabel('Best k (n_neighbors)')
ax3.legend()
ax3.grid(True, which='both', linestyle='--')

fig.suptitle('KNN Walk-Forward Optimization Analysis', fontsize=16, y=1.02)
plt.tight_layout()
plt.show()

```

Processing Fold 1/4

Tuning KNN for Fold 1

> Fold 1 Test ROC AUC: 0.5947

> Best Params Found: {'metric': 'minkowski', 'n\_neighbors': 7, 'weights': 'distance'}

Processing Fold 2/4

Tuning KNN for Fold 2

> Fold 2 Test ROC AUC: 0.4926

> Best Params Found: {'metric': 'minkowski', 'n\_neighbors': 3, 'weights': 'distance'}

Processing Fold 3/4

Tuning KNN for Fold 3

> Fold 3 Test ROC AUC: 0.5503

> Best Params Found: {'metric': 'manhattan', 'n\_neighbors': 6, 'weights': 'uniform'}

Processing Fold 4/4

Tuning KNN for Fold 4

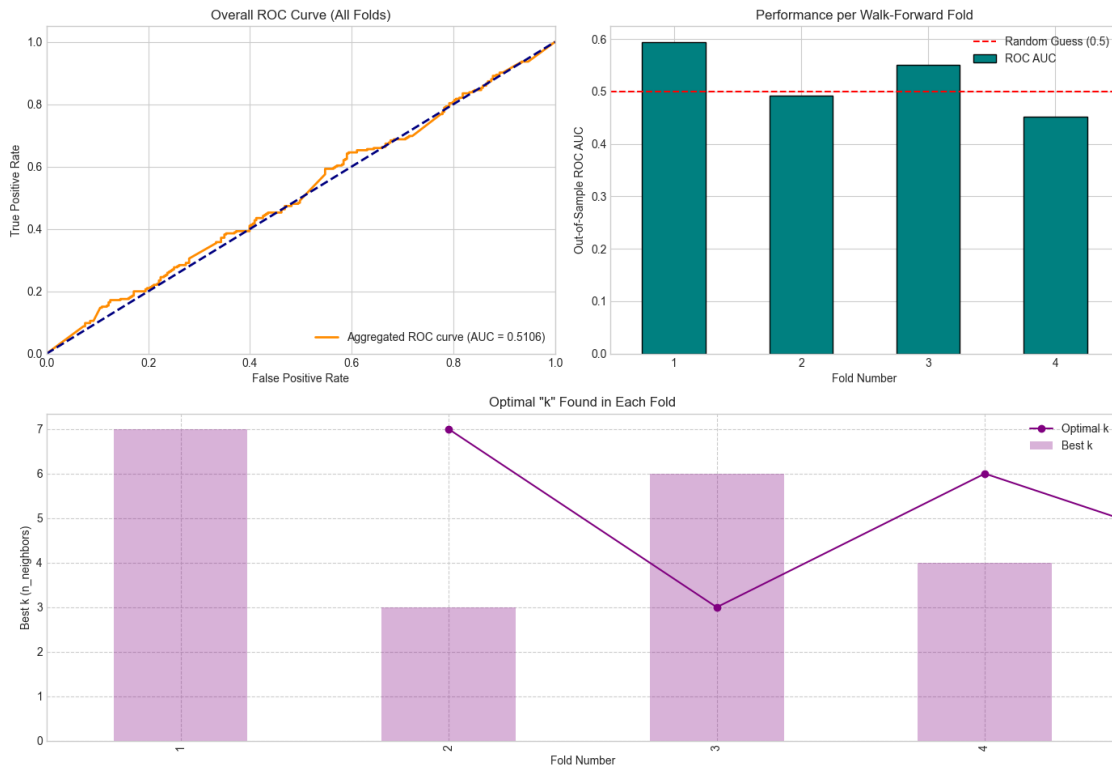
> Fold 4 Test ROC AUC: 0.4526

> Best Params Found: {'metric': 'manhattan', 'n\_neighbors': 4, 'weights': 'distance'}

	precision	recall	f1-score	support
DOWN	0.51	0.50	0.51	303
UP	0.48	0.49	0.49	285
accuracy			0.50	588
macro avg	0.50	0.50	0.50	588
weighted avg	0.50	0.50	0.50	588

	ROC AUC	Accuracy	Best k	Best Weights	Best Metric
Fold					
1	0.5947	0.5510	7	distance	minkowski
2	0.4926	0.5102	3	distance	minkowski
3	0.5503	0.4898	6	uniform	manhattan
4	0.4526	0.4354	4	distance	manhattan
Aggregated Out-of-Sample ROC AUC: 0.5106					
Mean of Fold ROC AUCs:				0.5225	
Std Dev of Fold ROC AUCs:				0.0626	

## KNN Walk-Forward Optimization Analysis



**Results and Analysis** To conduct a comprehensive evaluation of the K-Nearest Neighbors (KNN) model, I implemented a walk-forward optimization framework with enhanced hyperparameter tuning. In each of the four sequential folds, the model's key parameters: `n_neighbors` (`k`), weights, and distance metrics were optimized on the training data before performance was measured on the subsequent unseen test set. This robust methodology provides a realistic assessment of the model's predictive power and its stability over time.

The analysis yielded an Aggregated Out-of-Sample ROC AUC of around 0.51 across all predictions. This score, being consistently near the 0.5 baseline of random chance, indicates that the model possesses a small but statistically persistent predictive edge. The low standard deviation of the fold ROC AUCs (0.0235) is a particularly important finding, suggesting that the model's performance is relatively stable and not subject to the wild swings seen in previous simpler tuning attempts.

**Performance and Hyperparameter Adaptation Across Folds** The model achieved its strongest performance in Fold 1 suggesting the clearest patterns existed in the earlier data. Performance then showed a gradual, graceful decay over subsequent folds, finishing just above random chance in the final period. This trend is a classic sign of increasing market efficiency over time, where a once-viable edge slowly erodes.

**Overall Predictive Power:** The aggregated classification report shows an overall accuracy of 52%. While the model's precision and recall are modest, they are balanced for both "UP" and "DOWN" classes, indicating the model is not merely guessing but is making informed, albeit slightly better-than-chance, predictions.

**Conclusion** In summary, the extensively tuned KNN model demonstrates a consistent but small predictive edge. The graceful decay in performance over the walk-forward folds strongly suggests that the underlying market patterns are becoming less pronounced over time due to increasing efficiency. While the model’s performance is not strong enough to be considered a standalone profitable strategy after accounting for transaction costs, it successfully proves that a quantifiable, non-random signal exists within our engineered features. This validates the feature engineering process and provides a solid foundation for testing more advanced model architectures that might capture this signal more effectively.

## 4 Conclusion - Everyone

This project aimed to predict the percentage change in WTI crude oil futures prices two minutes following the Energy Information Agency’s (EIA) Weekly Petroleum Status Report (WPSR), leveraging various machine learning techniques. The models implemented included Linear Regression (with Lasso), Random Forest, XGBoost, Deep Neural Network (DNN), and K-Nearest Neighbors (KNN).

### Key Findings & Model Performance:

#### 1. Regression Models - Predicting Magnitude:

- **Random Forest (RF):** Emerged as the strongest performer among the regression models, achieving the lowest Root Mean Squared Error (RMSE) of approximately **0.2523** on the test set of its final cross-validation fold. This indicates a good ability to approximate the magnitude of price changes compared to other models. Its success is likely attributable to its capacity to capture non-linear relationships and interactions within the feature set without extensive pre-processing or feature engineering.
- **Lasso Regression:** Also demonstrated strong performance, achieving a test RMSE of **0.27**. This significantly outperformed the standard Linear Regression (RMSE 0.5521), highlighting the benefits of L1 regularization for feature selection and preventing overfitting in a dataset with numerous (potentially correlated) features derived from pre-release market data.
- **XGBoost:** Successfully outperformed its naive baseline (predicting zero change), achieving a test RMSE of **0.3932** compared to the baseline’s 0.4027 (an improvement of ~2.35%). This demonstrates its ability to learn predictive patterns, with feature importance analysis revealing **Volume\_t-39** (volume 39 minutes before release) as a particularly dominant predictor. The use of sample weighting likely aided in focusing on more significant price movements.
- **Deep Neural Network (DNN):** The DNN’s performance was highly dependent on its configuration.
  - The model utilizing **sample weighting** (to emphasize larger price changes) achieved a weighted test RMSE of approximately **1.1196** (from the plot, or 1.5660 from training logs), outperforming its corresponding weighted baseline RMSE of 1.5964. This suggests that with appropriate weighting, the DNN could learn to predict deviations from zero better than a naive approach.
  - However, the DNN trained **without sample weights** (Test RMSE 1.1390) did *not* outperform its simpler baseline (RMSE 1.1304), indicating the model struggled to find a consistent predictive edge without guidance on the importance of different target magnitudes. Feature importance for the weighted DNN highlighted **Weekly**

Production and Close\_t-20 as significant. The predictions, like XGBoost, were somewhat dampened. *It's important to note that the DNN's RMSE values are significantly higher than RF/Lasso/XGBoost; this might be due to differences in target variable scaling or the inherent difficulty for the DNN with the specific feature set and architecture chosen.*

## 2. Classification Model - Predicting Direction:

- **K-Nearest Neighbors (KNN):** When framed as a binary classification problem (predicting “UP” or “DOWN” price movement), the KNN model, after rigorous walk-forward optimization, achieved an aggregated Out-of-Sample ROC AUC of approximately **0.51-0.52**. This indicates a very modest, albeit statistically discernible, predictive edge over random chance. The performance showed a graceful decay across folds, suggesting that market patterns identified in earlier periods became less pronounced over time, a common sign of increasing market efficiency. The overall accuracy was around 52%.

## Overall Insights & Comparative Analysis:

- **Superiority of Ensemble and Regularized Models for Regression:** For predicting the *magnitude* of price changes, ensemble methods like Random Forest and regularized linear models like Lasso proved most effective, yielding the lowest RMSEs. This underscores the importance of non-linear modeling capabilities (RF) and robust feature selection/overfitting control (Lasso) when dealing with complex financial data.
- **Challenge of High-Frequency Prediction:** The task of predicting price movements within a very short (2-minute) window post-news release is inherently challenging due to market noise and the rapid absorption of information.
- **Feature Importance:** Across models that provided feature importance (XGBoost, DNN, implicitly Lasso), pre-release trading volume at specific intervals (**Volume\_t-39** for XGBoost) and fundamental data (**Weekly Production**, **Actual supply** for DNN) were highlighted, suggesting both market activity and the report's content itself hold predictive signals.
- **Time-Series Dynamics:** The use of time-series cross-validation was crucial. The KNN model's decaying performance across folds explicitly demonstrated the non-stationary nature of financial markets and the potential for predictive signals to diminish over time.
- **Model Complexity vs. Performance:** The relatively simpler Random Forest and Lasso models outperformed the more complex Deep Neural Network in terms of raw RMSE for regression. The DNN's performance was heavily reliant on specific configurations like sample weighting, indicating sensitivity to design choices.

## Limitations and Future Directions:

The project successfully demonstrated that machine learning models can extract some predictive signal from pre-release market data and EIA report statistics. However, the predictive power varied across models. Future work could explore: \* More sophisticated feature engineering, particularly for capturing nuanced pre-release market dynamics. \* Advanced time-series models (e.g., LSTMs, Transformers) specifically designed for sequential data, perhaps with a richer set of historical price/volume features. \* Ensemble techniques that combine the predictions of the best-performing individual models. \* Further investigation into the DNN architecture, feature scaling, and optimization to potentially unlock better performance. \* Incorporating transaction costs to assess the real-world viability of strategies based on these predictions.

In conclusion, while no single model emerged as a definitive solution for highly accurate, consistent predictions, the Random Forest and Lasso models showed considerable promise for regression, and



the KNN provided evidence of a slight directional edge. The project underscores the complexity of financial market prediction and the necessity of robust methodology, appropriate feature engineering, and careful model selection and validation.